# Interactive High-performance (HPC) Computing

**COMPLECS Team**

https://bit.ly/COMPLECS

https://github.com/sdsc-complecs/interactive-computing/

SDSC
SAN DIEGO SUPERCOMPUTER CENTER

UC San Diego

# About COMPLECS

COMPLECS (COMPrehensive Learning for end-users to Effectively utilize CyberinfraStructure) is a new SDSC program where training will cover non-programming skills needed to effectively use supercomputers. Topics include parallel computing concepts, Linux tools and bash scripting, security, batch computing, how to get help, data management and interactive computing.

**COMPLECS is supported by NSF award 2320934.**

# Outline

- Defining Interactive HPC
  - High-performance computing (HPC)
  - HPC batch computing
  - Interactive computing
- Accessing Interactive HPC Nodes
  - Launching nodes
  - Running GUIs using X11 forwarding
- Interactive Application Examples
  - Viewing Data: unix file ops (grep, awk, cat), gnuplot, NetCDF
  - Programming & Visualization Platforms: Matlab, R, Jupyter Notebooks
  - Gateways & Portals: simplify access to interactive apps
- Q&A

# You can run amazing jobs on supercomputers!

## Typically parallel or very large memory

*Pelagic fish communities Shapes (Jerome Guiet, UCLA)*

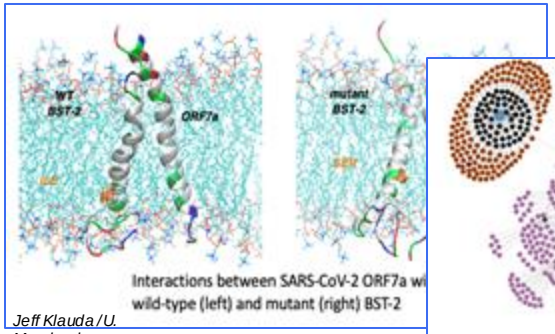


*Jeff Klauda / U. Maryland*

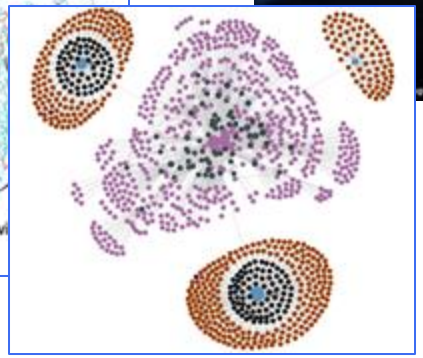

*Sample of Internet structure from CAIDA data (Mark Burgess, 12/16/21)*

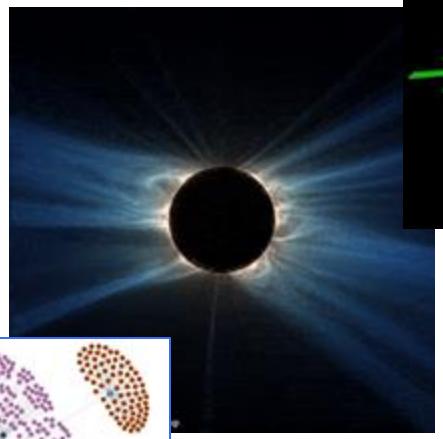

*Cooper Downs, Predictive Science Inc*

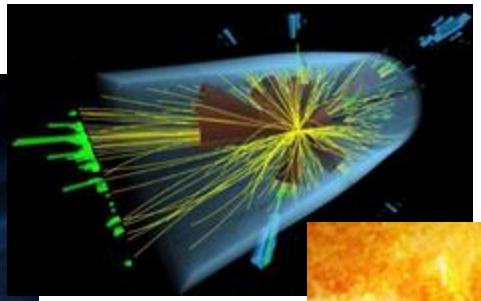

Compact Muon Solenoid (**CMS**)experiment at the LHC, CERN. *Image courtesy of CMS Collaboration; Mc Cauley, Thomas*

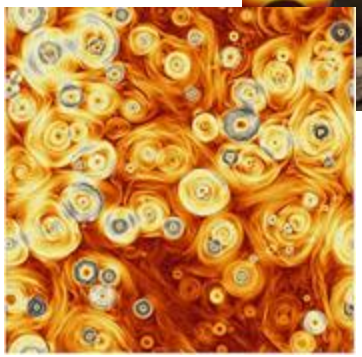

*Modeling the sun's corona, Alfred Mallet (UC Berkeley)*

*Model of Both Inner and Outer Solar System M. S.Clement (Carnegie Institution for Science)*

# However, you have to run your jobs on HPC clusters

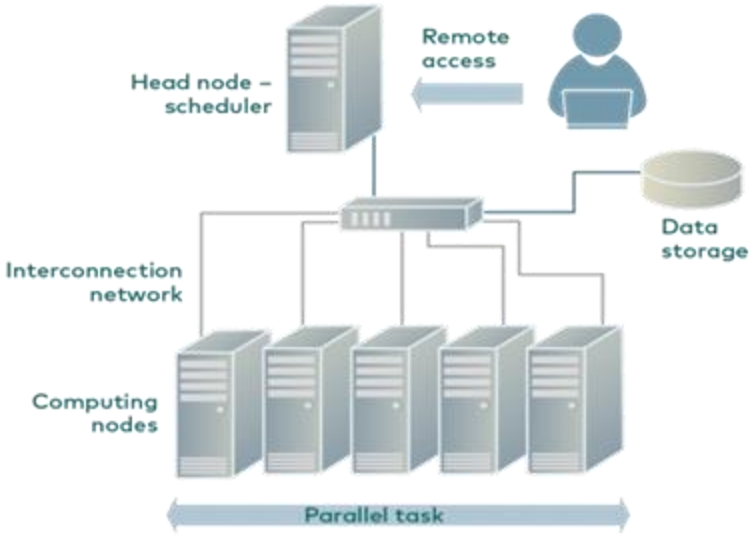


- Systems powerful but complex
- Jobs can be run from:
  - Command line interface (CLI)
  - The batch queue system
  - Application clients (PyTorch)
- Job scaling:
  - Parallel: 1 core to 1000's of nodes
- For large jobs, schedulers are needed to coordinate tasks

# HPC System Architecture: Expanse @ SDSC

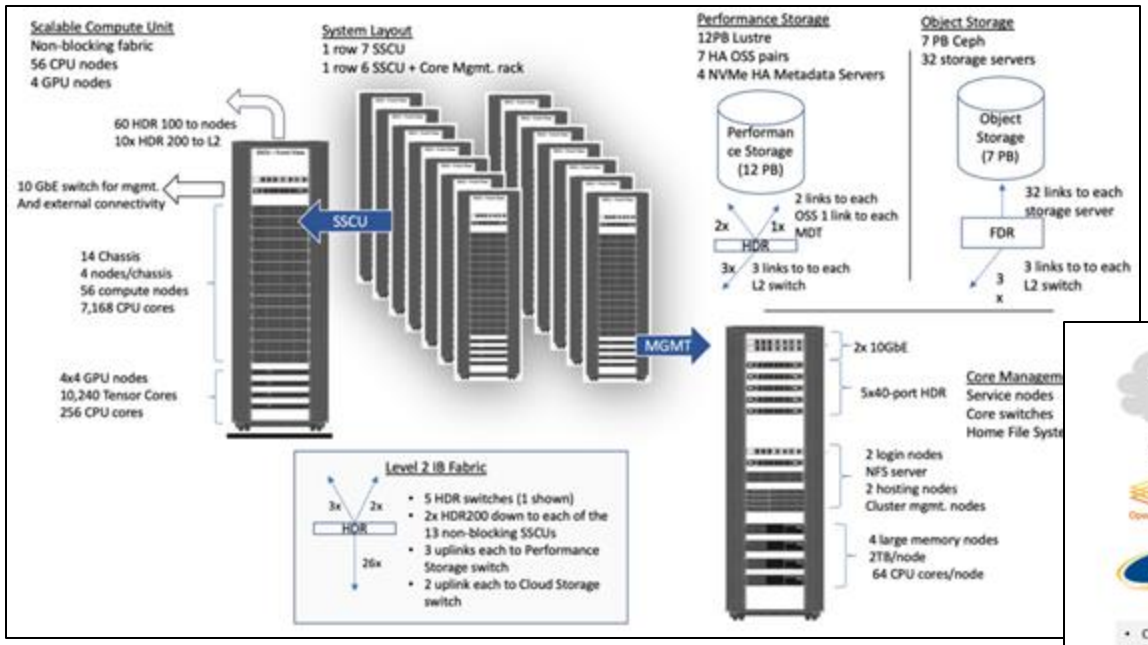


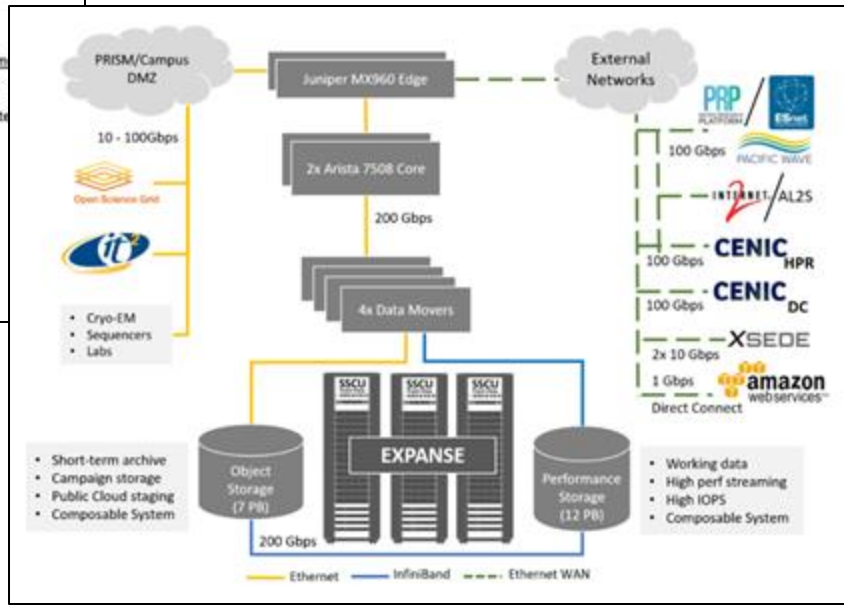


https://expanse.sdsc.edu
https://www.youtube.com/watch?v=uNZyg6X_t3s
https://access-ci.org

# HPC Jobs can be run as batch jobs (background) or interactively (real time)

Terminal shell – submit batch

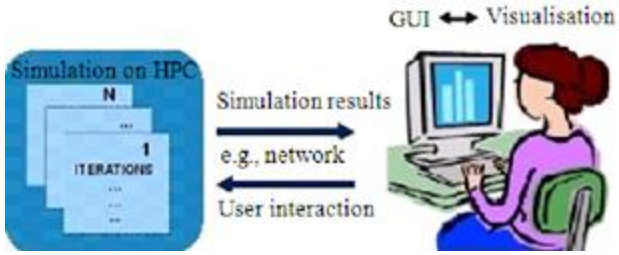


https://portal.expanse.edu

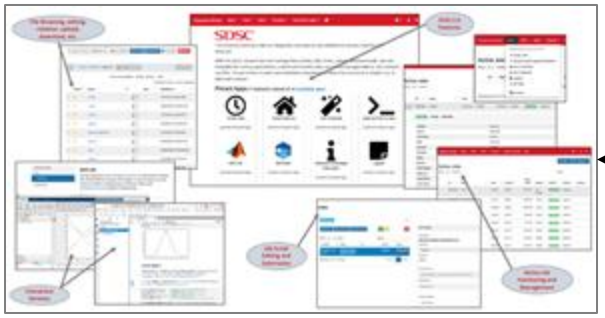

Interactive Distributed Computing with Jupyter

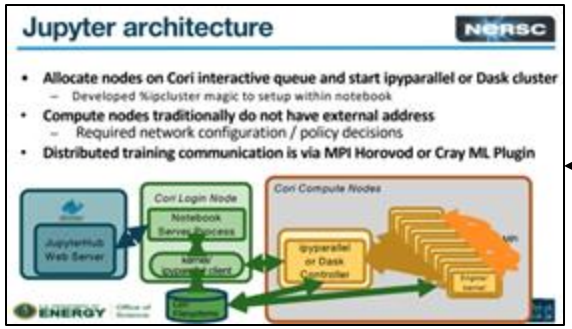


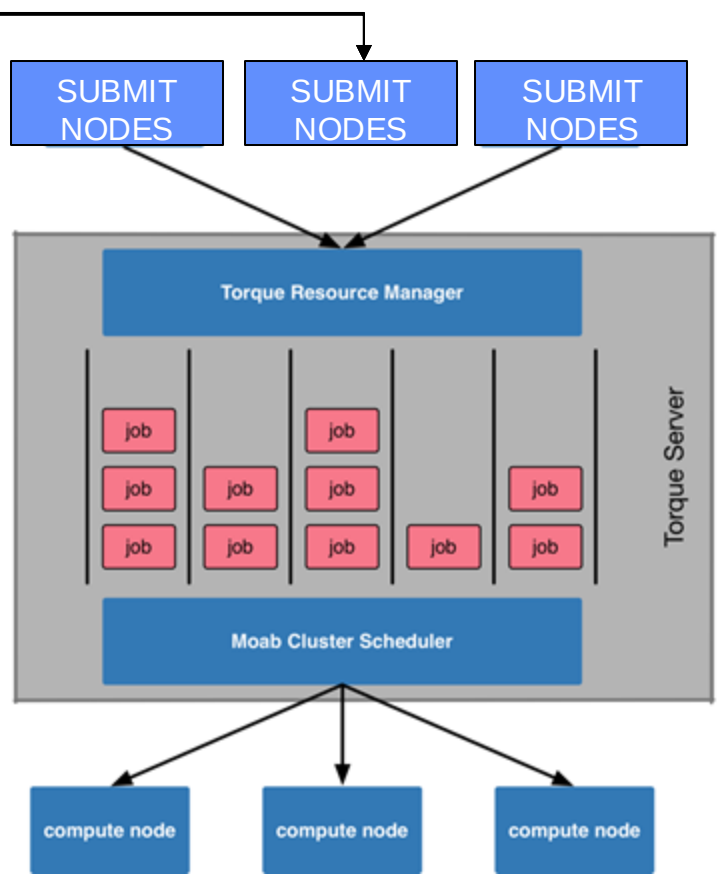


Src: https://hpc.dccn.nl/_images/torque_moab_arch.png

# Outline

- Defining Interactive HPC
  - High-performance computing (HPC)
  - HPC batch computing
  - Interactive computing
- Accessing Interactive HPC Nodes
  - Launching nodes
  - Running GUIs using X11 forwarding
- Interactive Application Examples
  - Viewing Data: unix file ops (grep, awk, cat), gnuplot, NetCDF
  - Programming & Visualization Platforms: Matlab, R, Jupyter Notebooks
  - Gateways & Portals: simplify access to interactive apps
- Q&A

# What is a batch job scheduler & how are they used?

- Any HPC (or HTC) system — usually a cluster of machines — needs a means of sharing computational resources fairly between users; without this, there would be anarchy.
- Batch-queueing systems — usually abbreviated to simply *batch systems* — are intended to do this.
- All batch systems have at least these features:
  - *a scheduler* for allocating resources (CPUs!) to jobs and for prioritizing jobs;
  - *one or more queues* to which jobs are submitted
  - Job partitions (or job queues): queue can be configured for a particular type of job (serial or parallel jobs; long or short jobs; or those requiring particularly high memory); members of a group or project).
- You need a basic understanding of how batch jobs work in order to run interactive jobs on HPC Clusters (ironic?)

# Batch Scheduler: main goals

- Minimize time between job submission and completion:
  - No job should stay in queue for extensive periods of time.
- Optimize CPU utilization:
  - Algorithms focus on minimizing CPU idle times.
- Maximize job throughput:
  - Manage as many jobs per time unit as possible.
- Support running jobs automatically in the background
- Finalize jobs: ensure job data and results are recorded & stored where you want them.

# Simple “Batch Scheduler” Architecture

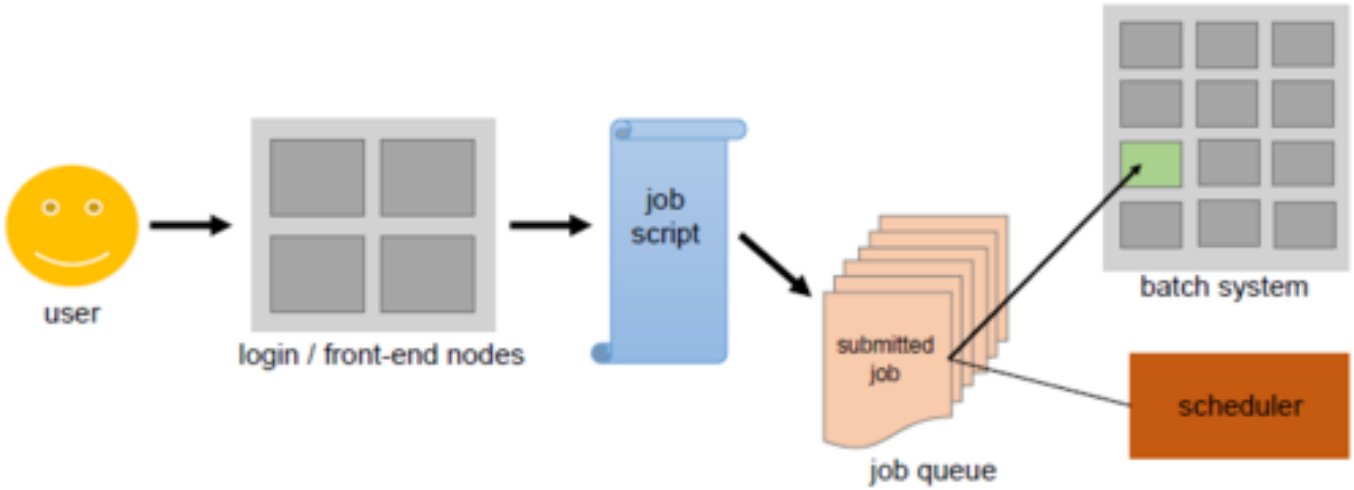


- Batch scheduler: software that implements a *batch system* on a cluster.
- Users do not run calculations interactively -- instead they submit *non-interactive batch jobs* to the *scheduler*.
- All work about the same: some are open source; some cost money; some are very expensive.

https://hpc-wiki.info/hpc/Scheduling_Basics

# Batch jobs: Slurm resource manager

Simple Linux Utility for Resource Management

- Open Source, runs on many systems
- “Glue” for parallel computer to schedule and execute jobs
- Role: Allocate resources within a cluster
  - Nodes (unique IP address)
  - Interconnect/switches
  - Generic resources (e.g. GPUs)
  - Launch and otherwise manage jobs

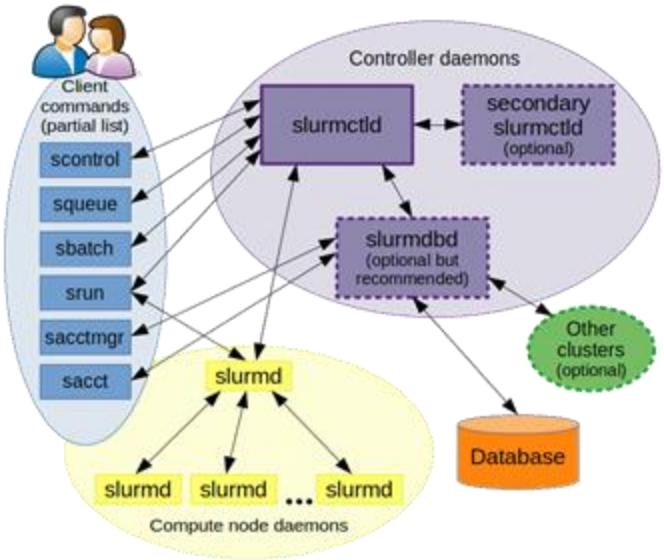


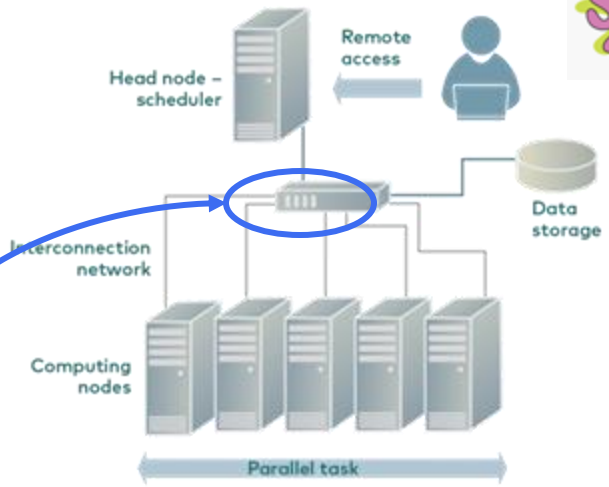


Functionality:

- Prioritize queue(s) of jobs;
- decide when and where to start jobs;
- terminate job when done;
- Appropriate resources;
- manage accounts for jobs

https://slurm.schedmd.com/sbatch.html

# How a scheduler schedules jobs

- Simple example:
  - 6-node system
  - Users submit 9 jobs.
  - Jobs 2 and 9 are multi-node
- Scheduler places the jobs in the queue and then onto the available nodes as they open up.
- Many parameters affect scheduling:
  - number of jobs submitted
  - required runtime
  - required number of cores
  - Required number of nodes
  - required main memory
  - CPU vs GPU
  - libraries, etc.

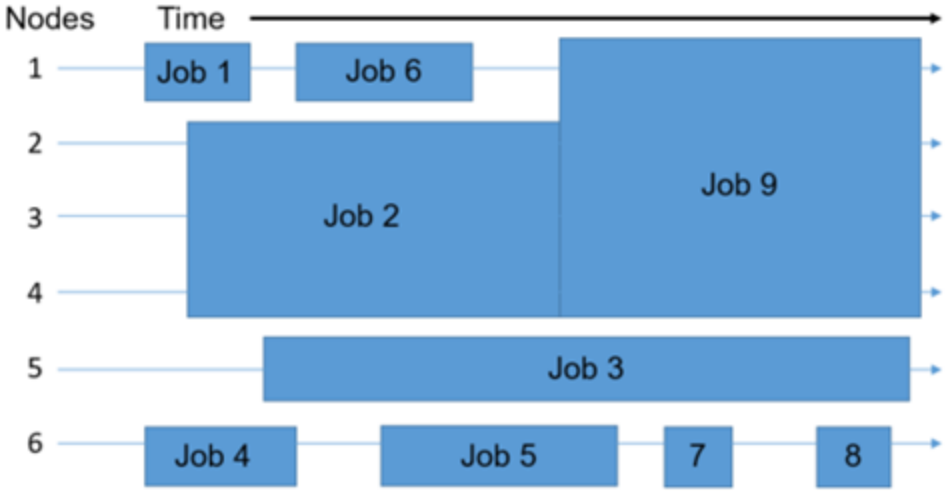


**Scheduler needs to play kind of "multidimensional tetris" to fill the cluster's nodes evenly and efficiently.**

https://hpc-wiki.info/hpc/Scheduling_Basics

# Outline

- Defining Interactive HPC
  - High-performance computing (HPC)
  - HPC batch computing
  - Interactive computing
- Accessing Interactive HPC Nodes
  - Launching nodes
  - Running GUIs using X11 forwarding
- Interactive Application Examples
  - Viewing Data: unix file ops (grep, awk, cat), gnuplot, NetCDF
  - Programming & Visualization Platforms: Matlab, R, Jupyter Notebooks
  - Gateways & Portals: simplify access to interactive apps
- Q&A

# What is Interactive HPC Computing

- In **computer** science, **interactive computing** refers to software which accepts input from the user as it runs.
  - **Interactive** software includes commonly used programs, such as word processors or spreadsheet applications.
- **Interactive HPC computing** involves *real-time* user inputs to perform tasks on a set of compute node(s) including:
  - Code development, real-time data exploration, and visualizations.
  - Used when applications have large data sets or are too large to download to local device, software is difficult install, etc.
  - User inputs come via command line interface or application GUI (Jupyter Notebooks, Matlab, R-studio).
  - Actions performed on remote compute nodes as a result of user input or program out.

# Interactive HPC Computing- Motivation

- Need more memory: Your jobs no longer fit onto the CPU/system you have been using:
  - That is why Expanse has 768 nodes (128 cores per node), with 256 GB DDR memory on each node
  - My MacBook Pro: 1 node, 8 cores, 16 GB DDR
- Too much data: your application needs more room:
  - Expanse has 1TB NVME/node, and 12 PB file system.
  - My MBP: 500 GB, no NVME
- Your network is too slow:
  - Expanse has connections to ~ 150GB/sec (or faster) networks
  - My MBP: 300 Mbps download; 11 Mbps upload
  - Bioinfo lab, running analysis on PC: FastQ dataset ~ 500 MB: would need 360+ seconds to upload 1 run.

# Real-time User Interactions

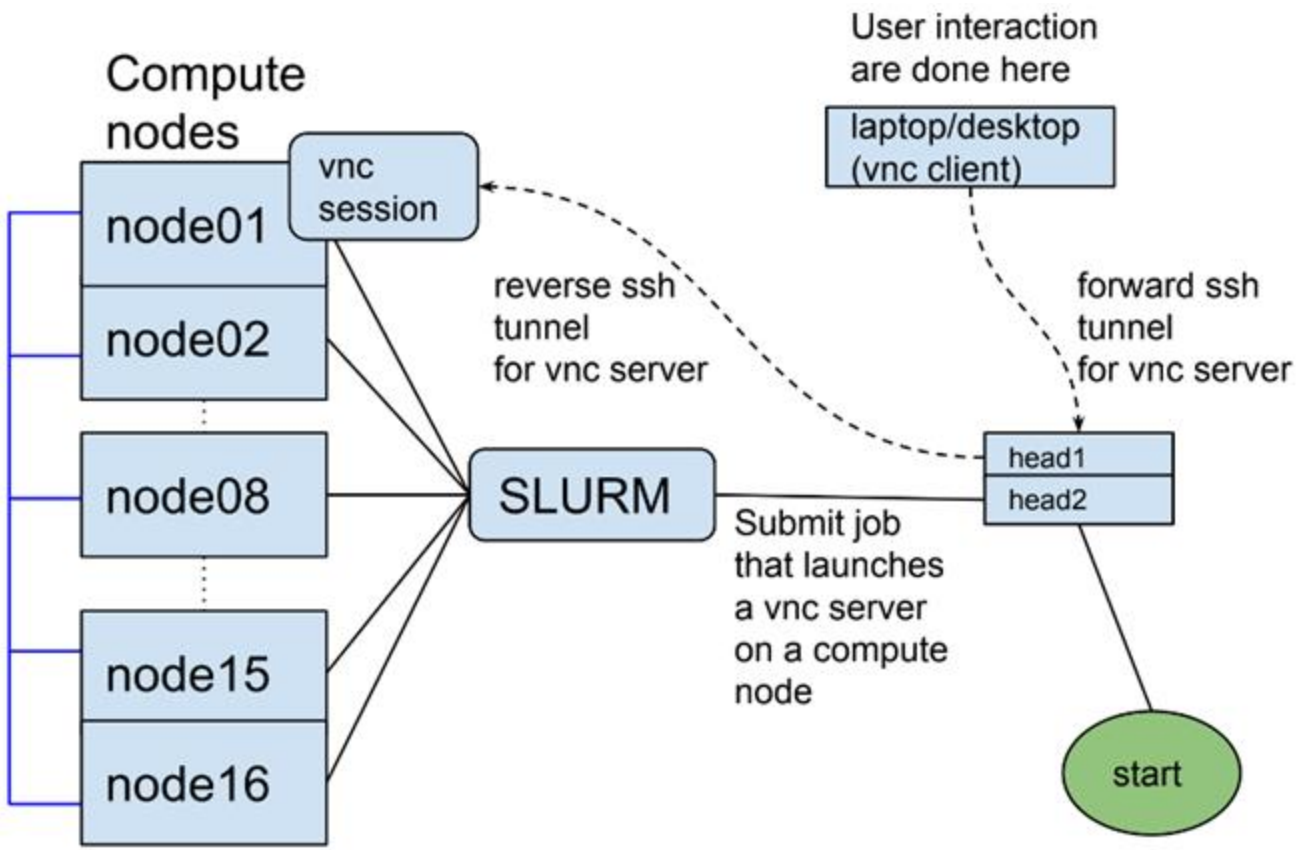


- *Real-time:* user inputs and performs tasks on a set of compute nodes
- Examples: Juptyer notebooks; ParaView visualization tools

# Interactive HPC Methods & Applications

- Interact with data after job is done:
  - Unix: query file info, location, output, grep, awk, sed
  - Cat the file contents from batch job or raw data
  - NetCDF data browser
- Plot results:
  - From within the code/model using libraries
  - Command-line driven graphing utility : Gnuplot
- Data visualization apps
  - NetCDF, HPF, TAU, ParaView
  - other
- Data Analysis Platforms: Matlab, R

# Interactive HPC Computing

https://jupyter.org/

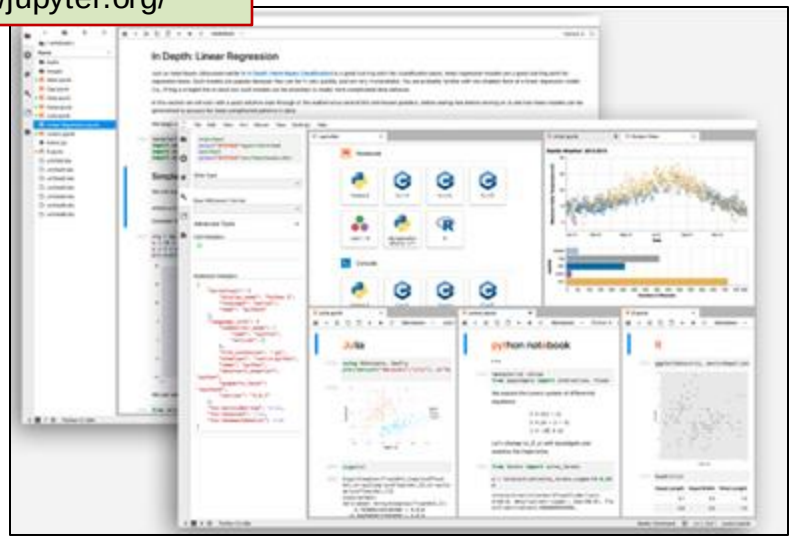

Parallel Matlab (AWS)

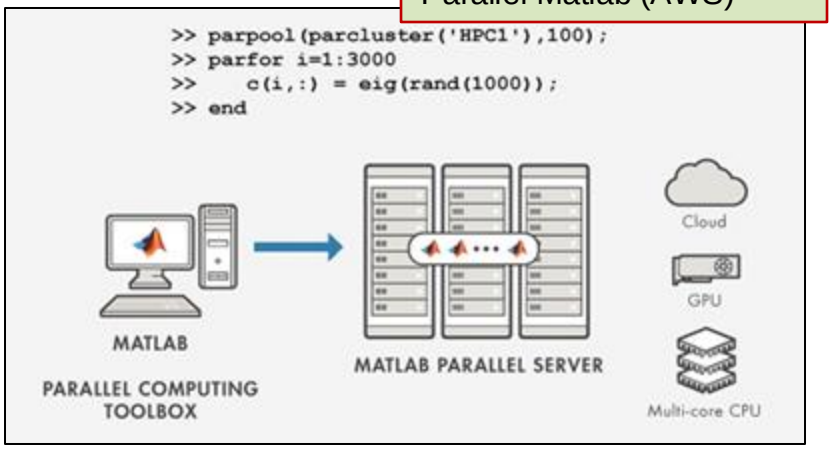


https://azuremarketplace.microsoft.com/en-us/marketplace/apps/mathworks-inc.matlab-parallel-server-listing?tab=Overview

Interactive Distributed Computing with Jupyter (NERSC)

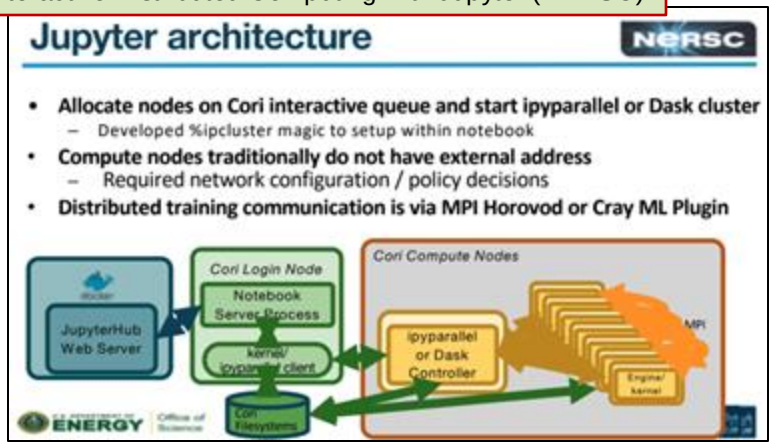


https://drive.google.com/file/d/1-OFjrk1q3L1d3uakr2xkozrPn2c2VZpZ/view

Paraview (AWS)

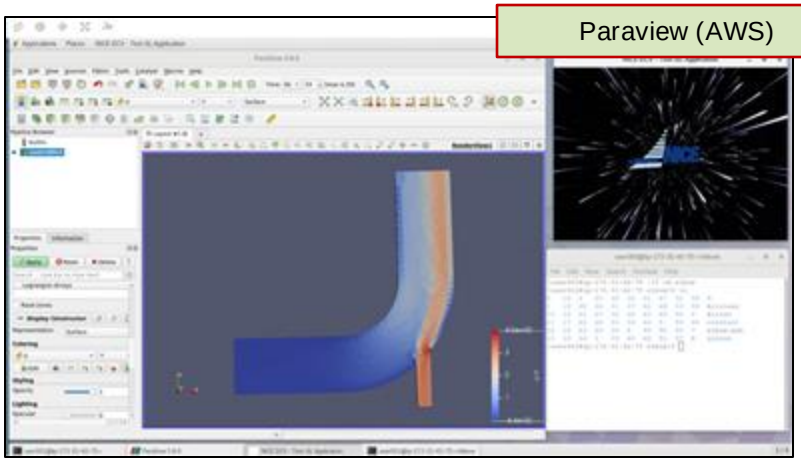

https://aws.amazon.com/blogs/compute/how-to-run-3d-interactive-applications-with-nice-dcv-in-aws-batch/

# Expanse User Portal

https://portal.expanse.sdsc.edu

- Provides Web based access to interactive applications including Jupyter Notebooks & Jupyter Lab, Matlab, Rstudio.
- Access using XSEDE portal account

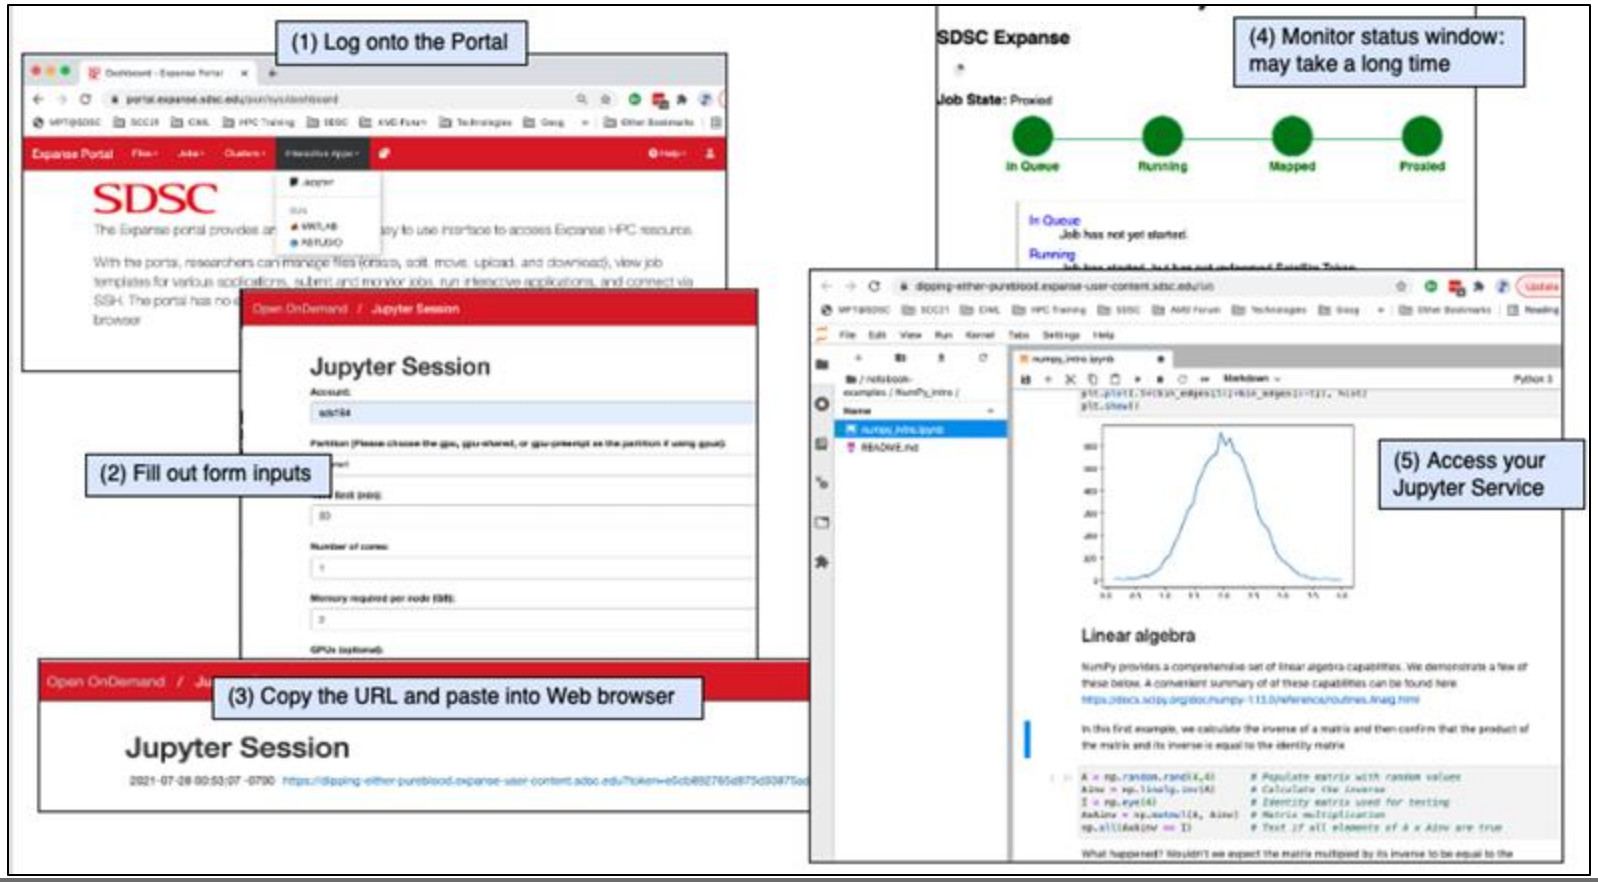

# Note: Interactive Services have a Key Vulnerability: They Provide Access to HPC File Systems, often over HTTP

SDSC Interactive Services Policy:

- Portals, JupyterHub, and other services cannot be mounted directly to disk (must be on VM or external):
  - Many use root in vulnerable ways
  - If a user launches Jupyter Lab or Notebooks, the jobs will be killed.
- Applications cannot run on login nodes.

SDSC recommendation:

- Use secure connections: when you choose insecure connections your account is vulnerable to hacking.
- Use portal.expanse.sdsc.edu

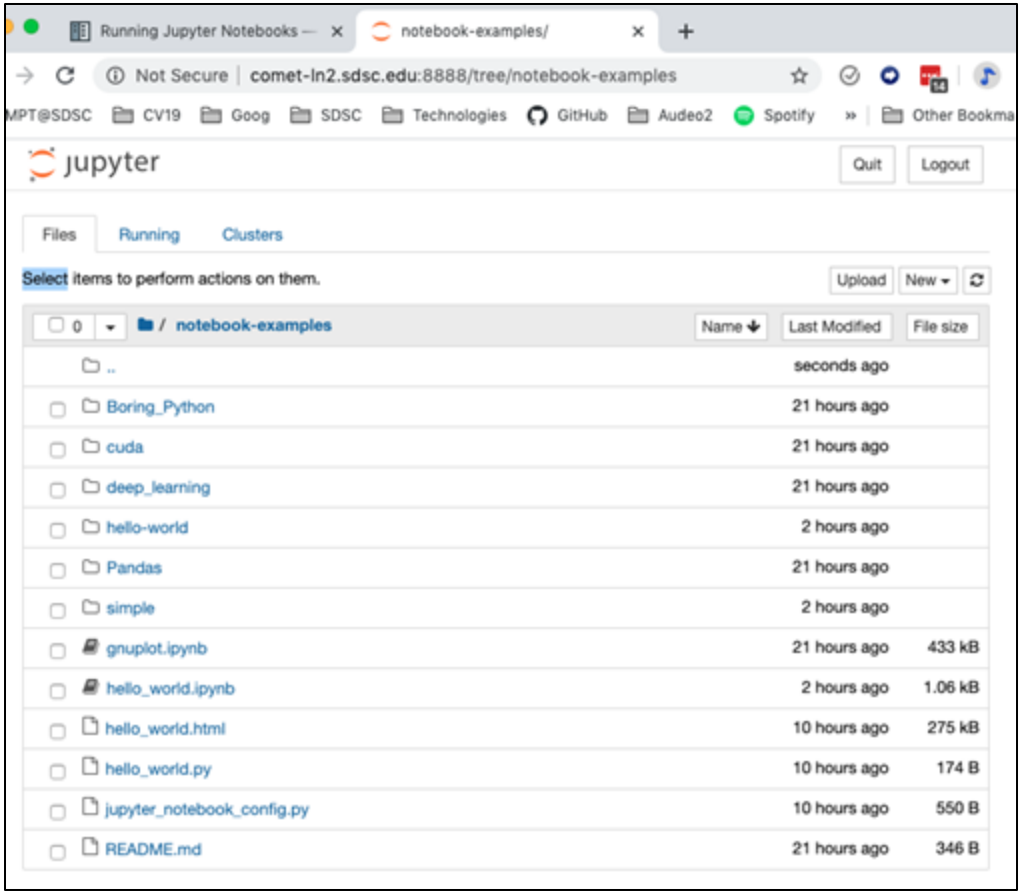

# Outline

- Defining Interactive HPC
  - High-performance computing (HPC)
  - HPC batch computing
  - Interactive computing
- Accessing Interactive HPC Nodes
  - Launching nodes
  - Running GUIs using X11 forwarding
- Interactive Application Examples
  - Viewing Data: unix file ops (grep, awk, cat), gnuplot, NetCDF
  - Programming & Visualization Platforms: Matlab, R, Jupyter Notebooks
  - Gateways & Portals: simplify access to interactive apps
- Q&A

# When would you need to use an Interactive node?

- The application is interactive (Matlab, Notebooks)
- Need to stage large amounts of data
- AI applications (Pytorch, Tensorflow)
- Need to compile a very large, complex application that may take a long time.
  - Java code, 1 million lines, > 30 minutes on DELL 5511 laptop

# Accessing Interactive Compute Nodes on Expanse

- Connect to HPC system (e.g. Expanse) via terminal using SSH  secure connections
- Use the *srun* command to obtain nodes for 'live,' command line interactive access:

| | |
|---|---|
| **CPU** | srun --partition=gpu-debug --pty --account=use300 --ntasks-per-node=10 --nodes=1 --mem=96G --gpus=1 -t 00:30:00 --wait=0 --export=ALL /bin/bash |
| **GPU** | srun --partition=gpu-debug --pty --account=use300 --ntasks-per-node=10 --nodes=1 --mem=96G --gpus=1 -t 00:30:00 --wait=0 --export=ALL /bin/bash |

(Tested 04/17/2024)

# Using An Interactive CPU node

```
[mthomas@login01 calc-prime]$ srun --partition=gpu-debug --pty --account=use300 --
ntasks-per-node=16 --nodes=1 --mem=96G --gpus=1 -t 00:30:00 --wait=0 --
export=ALL /bin/bash

srun: job 24457429 has been allocated resources
[mthomas@exp-9-55 calc-prime]$ module purge
[mthomas@exp-9-55 calc-prime]$ module load slurm
[mthomas@exp-9-55 calc-prime]$ module load cpu
[mthomas@exp-9-55 calc-prime]$ module load gcc/10.2.0
[mthomas@exp-9-55 calc-prime]$ module load openmpi/4.1.1
[mthomas@exp-9-55 calc-prime]$ mpirun -n 16 ./mpi_prime
06 August 2023 11:10:26 PM
PRIME_MPI     n_hi= 5000000      C/MPI version
An MPI example program to count the number of primes: # processes is 64
        N       Pi          Time
        1        0       0.013258
        2        1       0.001058
        4        2       0.000101
        8        4       0.000101
    [SNIP]
   131072    12251       0.110848
   262144    23000       0.410792
   524288    43390       1.527210
  1048576    82025       5.733612
  2097152   155611      21.725862
PRIME_MPI - Master process: Normal end of execution.
06 August 2023 11:12:26 PM
```

Request an interactive node for 30 minutes

- Exit interactive session when your work is done or you will be charged more CPU time.
- Beware of oversubscribing your job: don't ask for more cores than you have requested.
- Intel compiler allows this, but your performance will be degraded.

# Using Interactive GPU nodes

Request an interactive node for 30 minutes

```
[snip]
Last login: Fri Feb 18 12:58:32 2022 from 76.176.117.51
[username @login02 ~]$
[username @login02 ~]$ srun --partition=gpu-debug --pty --account=use300 --ntasks-per-node=10 --nodes=1 --
mem=96G --gpus=1 -t 00:30:00 --wait=0 --export=ALL /bin/bash
srun: job 9794018 queued and waiting for resources
srun: job 9794018 has been allocated resources
[mthomas@exp-14-57 ~]$
[mthomas@exp-14-57 ~]$ nvidia-smi
Fri Feb 18 13:04:19 2022
+-----------------------------------------------------------------------------+
| NVIDIA-SMI 460.32.03    Driver Version: 460.32.03    CUDA Version: 11.2     |
|-------------------------------+----------------------+----------------------+
| GPU  Name        Persistence-M| Bus-Id        Disp.A | Volatile Uncorr. ECC |
| Fan  Temp  Perf  Pwr:Usage/Cap|         Memory-Usage | GPU-Util  Compute M. |
|                               |                      |               MIG M. |
|===============================+======================+======================|
|   0  Tesla V100-SXM2...  On   | 00000000:86:00.0 Off |                    0 |
| N/A   34C    P0    41W / 300W |      0MiB / 32510MiB |      0%      Default |
|                               |                      |                  N/A |
+-------------------------------+----------------------+----------------------+

+-----------------------------------------------------------------------------+
| Processes:                                                                  |
|  GPU   GI   CI        PID   Type   Process name                  GPU Memory |
|        ID   ID                                                   Usage      |
|=============================================================================|
|  No running processes found                                                 |
+-----------------------------------------------------------------------------+
[username @exp-14-57 ~]$ exit
```

Verify you are on a GPU node

Exit when tasks are done

# Visualizating Apps: Use X11 Forwarding

- X Window System (X11): windowing system for bitmap displays, common on Unix-like operating systems.
  - MacOS uses Quartz
- X is an architecture-independent system for remote graphical user interfaces and input device capabilities

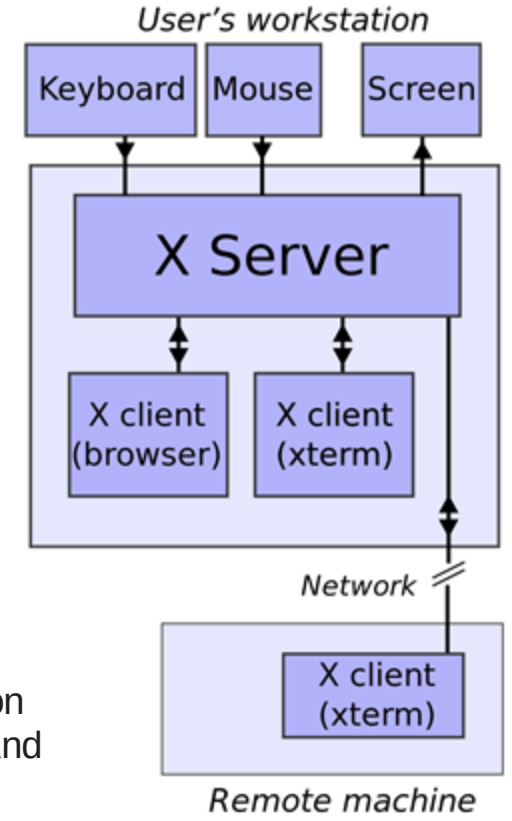


- X server receives input from a local keyboard and mouse and displays to a screen.
- Web browser and a terminal emulator run on the user's workstation
- A terminal emulator runs on a remote computer but is controlled and monitored from the user's machine

SRC: https://en.wikipedia.org/wiki/X_Window_System

# X11 Forwarding: connecting to Expanse

- On MacOS: Install X11 forwarding software:
  - MacOS: Xquartz
- Use the connection command:

  ssh -Y mthomas@login.expanse.sdsc.edu'

```
[mthomas@home]$ brew install --cask xquartz
Updating Homebrew...
Updated 2 taps (homebrew/core and homebrew/cask).
==> New Formulae
[SNIP]
installer: The upgrade was successful.
🍺  xquartz was successfully installed!
quantum:~ mthomas$ which xquartz
/opt/X11/bin/xquartz
```

Note: MacOS/Ventura has an X11 problem that is being debugged.

# X11 Forwarding to Expanse Interactive Node

- Use Xquartz windows
- Create two Expanse X11 login connections
  ssh -Y mthomas@login.expanse.sdsc.edu
- Window/Connection #1: request interactive node

```
[localhost]. ssh -Y mthomas@login.expanse.sdsc.edu
[mthomas@login01 ~] srun --partition=debug  --pty --account=use300 --nodes=1
--ntasks-per-node=4  --mem=8G -t 00:30:00 --wait=0 --export=ALL /bin/bash
srun: job 24457429 has been allocated resources
[mthomas@exp-9-55]$
```

- Window/Connection #2: connect to the interactive node; set up module environment and launch app

```
[localhost] ssh -Y mthomas@login.expanse.sdsc.edu
[mthomas@login01 ~] ssh -Y exp-9-55
Last login: Wed Dec  6 19:21:56 2023 from 10.21.0.19
[mthomas@exp-9-55 ~]$ module load  gpu/0.17.3b
[mthomas@exp-9-55 ~]$ module load matlab/2022b/nmbr5dd
[mthomas@exp-9-55 ~]$ matlab
```

- Note: MacOS X11 problem that is being debugged. Use the Expanse Portal

# gnuplot

- *Depends on X11, and your local host (MacOS XQuartz not working fully)*
- *gnuplot* is a command-driven interactive function plotting program. It can be used to plot functions and data points in both two- and three- dimensional plots in many different formats.
- Designed primarily for the visual display of scientific data.
- Can be run from the gnuplot app or from within your program so you can save visualizations of results.

cpu/0.15.4 gcc/10.2.0

```
[mthomas@exp-9-55]$ module load cpu/0.17.3b  gcc/10.2.0/npcyll4
gnuplot/5.4.2/mfinpvw
[mthomas@exp-9-55]$ which gnuplot
/cm/shared/apps/spack/cpu/opt/spack/linux-centos8-zen2/gcc-
10.2.0/gnuplot-5.2.8-uwugzxg4dgxaiciheiepgol67cw7m6yg/bin/gnuplot
[mthomas@login01 ~]$ gnuplot
        G N U P L O T
        Version 5.4 patchlevel 2    last modified 2021-06-01
        Copyright (C) 1986-1993, 1998, 2004, 2007-2021
        Thomas Williams, Colin Kelley and many others
        gnuplot home:     http://www.gnuplot.info
        faq, bugs, etc:   type "help FAQ"
        immediate help:   type "help"  (plot window: hit 'h')
Terminal type is now 'unknown'
gnuplot> set terminal xterm
Terminal type is now 'xterm'
gnuplot> plot sin(x)
```

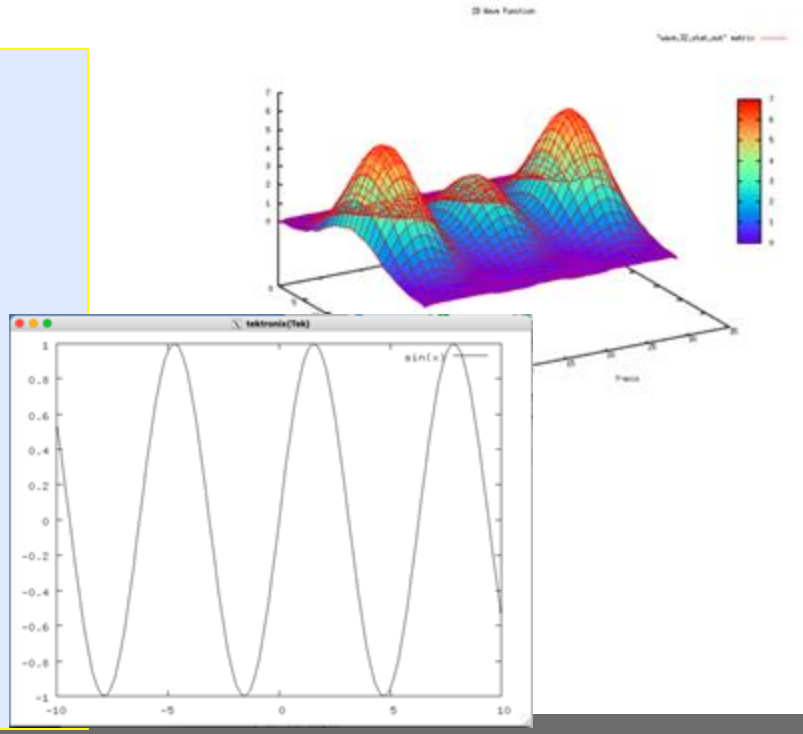

https://alvinalexander.com/technology/gnuplot-charts-graphs-examples/

# Outline

- Defining Interactive HPC
  - High-performance computing (HPC)
  - HPC batch computing
  - Interactive computing
- Accessing Interactive HPC Nodes
  - Launching nodes
  - Running GUIs using X11 forwarding
- Interactive Application Examples
  - Viewing Data: Unix file ops (grep, awk, cat), gnuplot, NetCDF
  - Programming & Visualization Platforms: Matlab, R, Jupyter Notebooks
  - Gateways & Portals: *simplify access* to interactive apps
- Q&A

# NetCDF Data Files

- NetCDF (Network Common Data Form) files are binary files containing both data and metadata about the data.
- API: software libraries and machine-independent data formats that support the creation, access, and sharing of array-oriented scientific data
- Similar to high-density format (HDF)
- NetCDF clients (*ncview*, *ncdump*) can be used to query and plot data in real-time

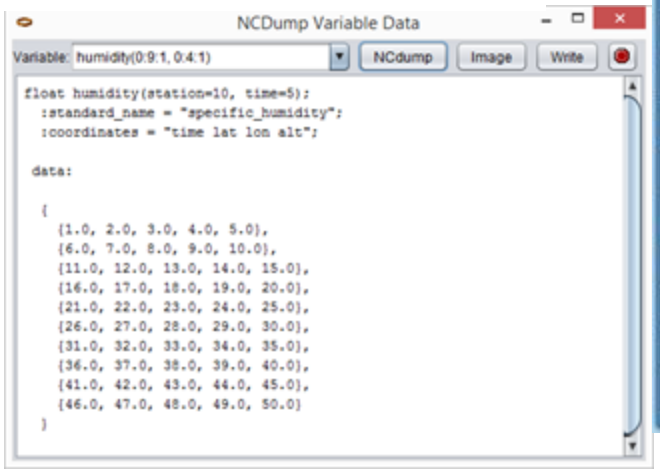


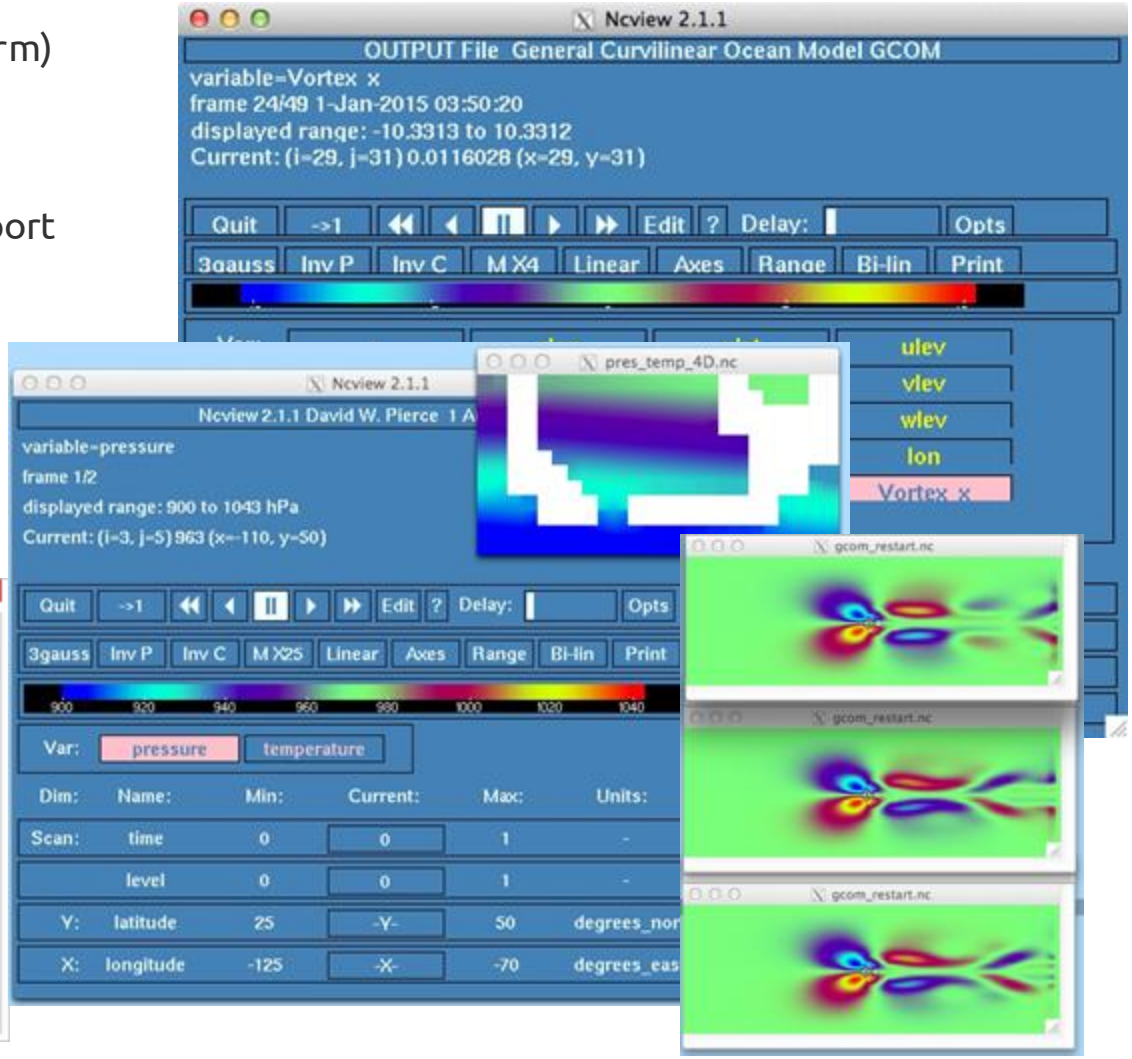


https://unidata.ucar.edu/netcdf

# Configuring NetCDF on Expanse

- Example below is for very simple, serial example provided by SciNet (Canada): see
  - Select the right module environment: it's a big software system
- Set up modules
- Put SciNet ENV vars into bash scripts
- Build the compile command (Makefile not easy to adapt)
- Example will generate a data.nc file
- Browse & query file contents

# NetCDF: Configure Shell Environment

- Add environment variables to your shell env:

```
[mthomas@login02 netcdf]$ cat ~/.bash_profile
[SNIP]
# netcdf library and include paths
SCINET_NETCDF_INC=/cm/shared/apps/spack/cpu/opt/spack/linux-centos8-zen2/intel-19.1.1.217/netcdf-c-4.7.4-62xm3fwd6soo2zqqbd4t7jza4ro5riyv/include
SCINET_NETCDF_LIB=/cm/shared/apps/spack/cpu/opt/spack/linux-centos8-zen2/intel-19.1.1.217/netcdf-c-4.7.4-62xm3fwd6soo2zqqbd4t7jza4ro5riyv/lib

#HDF5 library and include paths
SCINET_HDF5_INC=/cm/shared/apps/spack/cpu/opt/spack/linux-centos8-zen2/intel-19.1.1.217/hdf5-1.10.7-jsee7cinki6pcunimgfe3ovv6lz7gwnt/include
SCINET_HDF5_LIB=/cm/shared/apps/spack/cpu/opt/spack/linux-centos8-zen2/intel-19.1.1.217/hdf5-1.10.7-jsee7cinki6pcunimgfe3ovv6lz7gwnt/lib

#update path to binaries
PATH=$PATH:/cm/shared/apps/spack/cpu/opt/spack/linux-centos8-zen2/intel-19.1.1.217/netcdf-c-4.7.4-62xm3fwd6soo2zqqbd4t7jza4ro5riyv/bin:/cm/shared/apps/spack/cpu/opt/spack/linux-centos8-zen2/intel-19.1.1.217/hdf5-1.10.7-jsee7cinki6pcunimgfe3ovv6lz7gwnt/bin:/cm/shared/apps/spack/cpu/opt/spack/linux-centos8-zen2/intel-19.1.1.217/ncview-2.1.8-bm3tockq6ljgtxefovxia5hujnqmpdpn/bin
[SNIP]
```

# NetCDF: Configure  Modules

- Setup the right modules for the version you want to use:

```
[mthomas@login02 netcdf]$
[mthomas@login02 netcdf]$ module load cpu/0.15.4  intel/19.1.1.217  mvapich2/2.3.4 module load netcdf-c/4.7.4
[mthomas@login02 netcdf]$
[mthomas@login02 netcdf]$
[mthomas@login02 netcdf]$ which ncdump
/cm/shared/apps/spack/cpu/opt/spack/linux-centos8-zen2/intel-19.1.1.217/netcdf-c-4.7.4-
62xm3fwd6soo2zqqbd4t7jza4ro5riyv/bin/ncdump
[mthomas@login02 netcdf]$ which ncview
/cm/shared/apps/spack/cpu/opt/spack/linux-centos8-zen2/intel-19.1.1.217/ncview-2.1.8-
bm3tockq6ljgtxefovxia5hujnqmpdpn/bin/ncview
```

# Query NetCDF Data File Contents: ncdump

```
[mthomas@login02 netcdf]$ pwd
[mthomas@login02 netcdf]$ pwd
/home/mthomas/hpctr-examples/netcdf/parIO/netcdf
[mthomas@login02 netcdf]$ ls *.c
-rw-r--r-- 1 mthomas use300   8874 Oct  5  2010 2darray.c
-rwxr-xr-x 1 mthomas use300 304648 Dec  6 01:27 2darray-simple.c
-rw-r--r-- 1 mthomas use300  15903 Oct  6  2010 parallel2darray.c
-rw-r--r-- 1 mthomas use300   5865 Oct  5  2010 read2darray-simple.c
-rw-r--r-- 1 mthomas use300 240176mpicc -I/cm/shared/apps/spack/cpu/opt/spack/linux-centos8-zen2/intel-
19.1.1.217/netcdf-c-4.7.4-62xm3fwd6soo2zqqbd4t7jza4ro5riyv/include -L/cm/shared/apps/spack/cpu/opt/spack/linux-
centos8-zen2/intel-19.1.1.217/netcdf-c-4.7.4-62xm3fwd6soo2zqqbd4t7jza4ro5riyv/lib -lnetcdf -o 2darray 2darray.c
 Dec  6 01:27 data-simple.nc
[mthomas@login02 netcdf]$ ./2darray
[mthomas@login02 netcdf]$ ll *.nc
-rw-r--r-- 1 mthomas use300 241164 Dec  6 01:25 data.2darray.nc
-rw-r--r-- 1 mthomas use300 241164 Dec  6 21:27 data.nc
```

# Query NetCDF Data File Contents: ncdump

Tool for viewing output data file contents

```
[mthomas@login02 netcdf]$ ncdump -h
ncdump [-c|-h] [-v ...] [[-b|-f] [c|f]] [-l len] [-n name] [-p n[,n]] [-k] [-x] [-s] [-t|-i] [-g ...] [-w] [-Ln] file
  [-c]            Coordinate variable data and header information
  [-h]            Header information only, no data
  [-v var1[,...]]  Data for variable(s) <var1>,... only
  [-b [c|f]]      Brief annotations for C or Fortran indices in data
  [-f [c|f]]      Full annotations for C or Fortran indices in data
  [-l len]        Line length maximum in data section (default 80)
  [-n name]       Name for netCDF (default derived from file name)
  [-p n[,n]]      Display floating-point values with less precision
  [-k]            Output kind of netCDF file
  [-s]            Output special (virtual) attributes
  [-t]            Output time data as date-time strings
  [-i]            Output time data as date-time strings with ISO-8601 'T' separator
  [-g grp1[,...]]  Data and metadata for group(s) <grp1>,... only
  [-w]            With client-side caching of variables for DAP URLs
  [-x]            Output XML (NcML) instead of CDL
  [-Xp]           Unconditionally suppress output of the properties attribute
  [-Ln]           Set log level to n (>= 0); ignore if logging not enabled.
  file            Name of netCDF file (or URL if DAP access enabled)
```

# Read NetCDF Data File Contents: ncdump

***ncdump***: Tool for viewing output data file contents

```
[mthomas@login02 netcdf]$ ll *.nc
-rw-r--r-- 1 mthomas use300 241164 Dec  6 01:25 data.2darray.nc
-rw-r--r-- 1 mthomas use300 240176 Dec  6 01:27 data-simple.nc
[mthomas@login02 netcdf]$ ncdump -h data-simple.nc
netcdf data-simple {
dimensions:
X = 100 ;
Y = 100 ;
velocity\ component = 2 ;
variables:
double Density(X, Y) ;
double Velocity(velocity\ component, X, Y) ;
}
[mthomas@login02 netcdf]$ ncdump -x data-simple.nc
<?xml version="1.0" encoding="UTF-8"?>
<netcdf xmlns="http://www.unidata.ucar.edu/namespaces/netcdf/ncml-2.2" location="data-simple.nc">
  <dimension name="X" length="100" />
  <dimension name="Y" length="100" />
  <dimension name="velocity component" length="2" />
  <variable name="Density" shape="X Y" type="double">
  </variable>
  <variable name="Velocity" shape="velocity component X Y" type="double">
  </variable>
</netcdf>
```

# Read NetCDF Data File Contents: ncdump

```
[mthomas@login02 netcdf]$ ncdump -t  data-simple.nc | more
netcdf data-simple {
dimensions:
X = 100 ;
Y = 100 ;
velocity\ component = 2 ;
variables:
double Density(X, Y) ;
double Velocity(velocity\ component, X, Y) ;
data:
 Density =
  1.07326255555494, 1.07930091286413, 1.08569972700294, 1.09246679805167,
    1.09960871294202, 1.10713070597055, 1.11503651964736, 1.12332826729031,
    1.13200629886235, 1.14106907162359, 1.1505130272287, 1.16033247694148,
    1.17051949666066, 1.18106383345281, 1.19195282526666, 1.20317133545959,
    1.21470170369791, 1.22652371469942, 1.23861458616838, 1.25094897713016,
    1.26349901770561, 1.27623436117604, 1.28912225897801, 1.30212765903751,
    1.31521332760545, 1.3283399944956, 1.34146652135377, 1.35455009230782,
    1.3675464260648, 1.38041000823913, 1.39309434241745, 1.40555221819654,
[SNIP]
    -2.8468704354508, -2.69398173507181, -2.54522816466233,
    -2.40084389414749, -2.2610296301804, -2.12595326718743,
    -1.99575075552494, -1.87052716603845, -1.75035792940546,
    -1.63529022805598, -1.52534451818797, -1.42051615941371,
    -1.32077712986823, -1.22607780516019, -1.13634878032523, -1.05150271492394 ;
```

# Visualize NetCDF Data: ncview

```
[mthomas@login02 netcdf]$ ll *.nc
-rw-r--r-- 1 mthomas use300 241164 Dec  6 01:25 data.2darray.nc
-rw-r--r-- 1 mthomas use300 240176 Dec  6 01:27 data-simple.nc
[username@login02 netcdf] ncview data-simple.nc
```

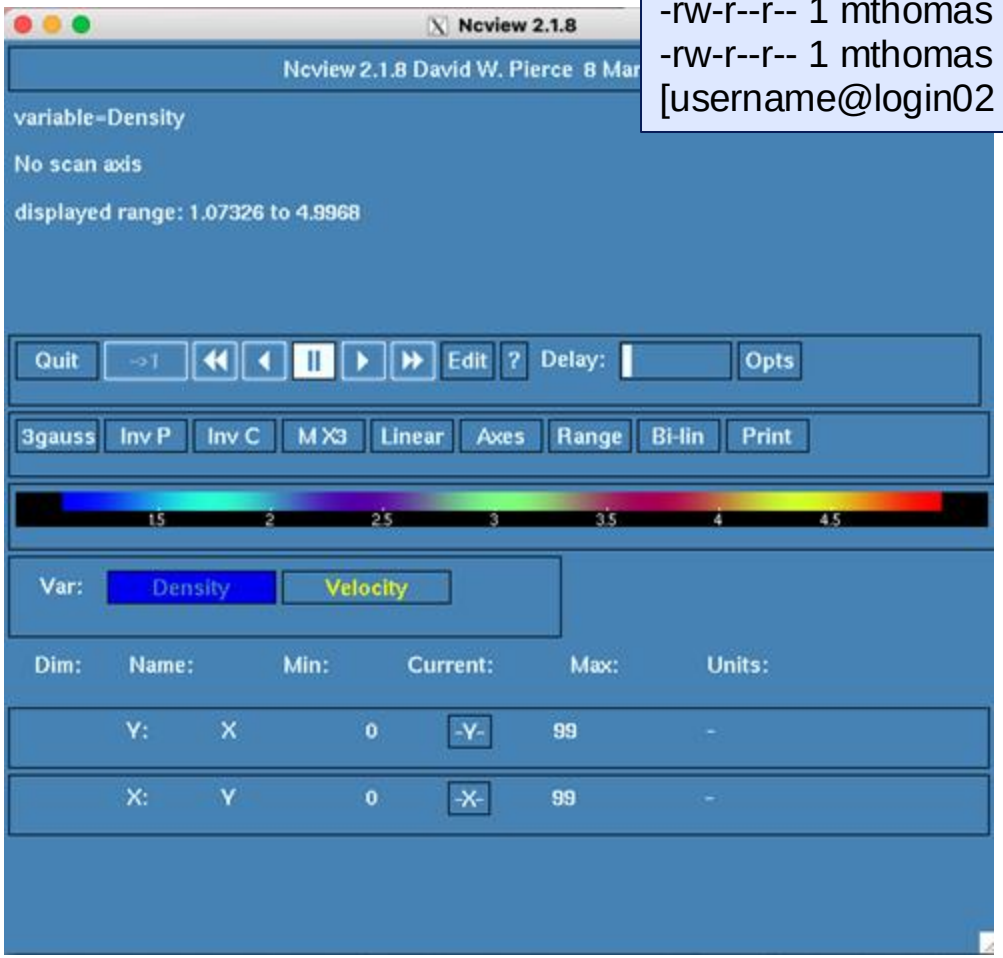


Viewing 2D density plot

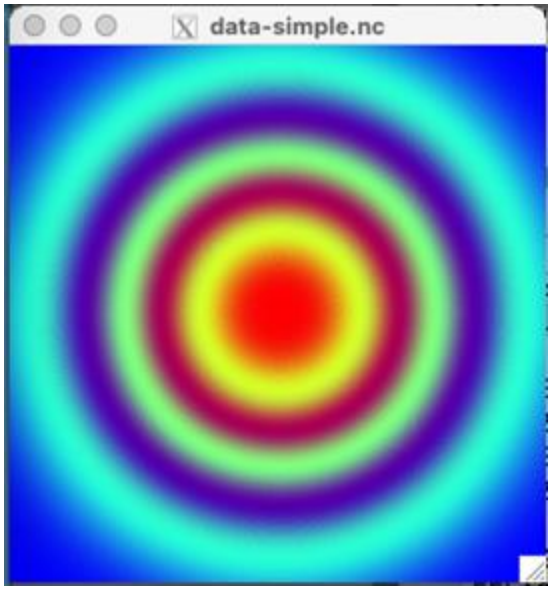

# Outline

- What is High-Performance Computing (HPC)?
- What is HPC batch computing
- Defining Interactive HPC Computing
- Accessing Interactive HPC Nodes
- Interactive Application Examples
  - Viewing Data: unix file operators (grep, awk, cat), gnuplot, NetCDF
  - Programming & Visualization Platforms: Matlab, R, Jupyter Notebooks (and NetCDF)
  - Gateways & Portals: simplify access to interactive apps

# Expanse Interactive Jobs: Running Matlab from Command Line

```
[mthomas@login01 ~]$  srun --partition=gpu-debug --pty --account=use300 --ntasks-per-node=10     --
nodes=1 --mem=96G --gpus=1 -t 00:30:00 --wait=0 --export=ALL /bin/bash
srun: job 14833549 queued and waiting for resources
srun: job 14833549 has been allocated resources
[mthomas@exp-9-55 ~]$ module load  gpu/0.17.3b
[mthomas@exp-9-55 ~]$ module load matlab/2022b/nmbr5dd
[mthomas@exp-9-55 ~]$ module list
Currently Loaded Modules:
  1) shared              3) sdsc/1.0       5) gpu/0.17.3b        (g)
  2) slurm/expanse/21.08.8   4) DefaultModules   6) matlab/2022b/nmbr5dd
```

Check that you have the right modules loaded

```
[mthomas@exp-9-55 ~]$ matlab
MATLAB is selecting SOFTWARE OPENGL
rendering.                                          < M A T L A B (R)>
  Copyright 1984-2022 The MathWorks, Inc.
     R2022a (9.12.0.1884302) 64-bit (glnxa64)
                February 16, 2022
 >> a = [1 3 5; 2 4 6; 7 8 10]
a =
     1     3     5
     2     4     6
     7     8    10
>> b=sin(a)
b =
    0.8415    0.1411   -0.9589
    0.9093   -0.7568   -0.2794
    0.6570    0.9894   -0.5440
>> plot(b)      -- nothing happens
>>
```

# Expanse Interactive Jobs: Running Matlab with GUI

To use a Graphical User Interface (GUI) as part of your interactive job, you will need to set up Xforwarding. Example below is for using XQuartz on a MAC. For examples for Windows, see: http://systems.eecs.tufts.edu/x11-forwarding/

***Step 1:*** *Log on to Expanse, using* ***–Y*** *option (trusted); create two connections; W2 should be the XQuartz terminal*

W1
```
[mthomas@home]$ ssh –Y mthomas@login.expanse.sdsc.edu
```

W2
```
[mthomas@home]$ ssh –Y mthomas@login.expanse.sdsc.edu
```

***Step 2: Window 1:*** *Request an interactive node, using* ***srun command***

*This window will not be used for anything else but holding the interactive node connection*

W1
```
[mthomas@login01 ~]$ srun --partition=debug  --pty --account=use300 --nodes=1 --ntasks-per-node=4  --mem=8G -t 00:30:00 --wait=0 --export=ALL /bin/bash
srun: job 26814162 queued and waiting for resources
srun: job 26814162 has been allocated resources
[mthomas@exp-9-55 ~]$
```

***Step 3: Window 2:***
***Connect to the interactive node:*** *atlab*

W2
```
[mthomas@login02 ~]$ ssh -Y exp-9-55
Last login: Wed Dec  6 19:23:47 2023 from 10.21.0.19
```

***Step 4: Window 2:***

*Setup Matlab module environment*
*Launch Matlab*

GUI : USE EXPANSE PORTAL

W2
```
[mthomas@exp-9-55 ~]$ module load  gpu/0.17.3b
[mthomas@exp-9-55 ~]$ module load matlab/2022b/nmbr5dd
[mthomas@exp-9-55 ~]$ module list
Currently Loaded Modules:
  1) shared     3) sdsc/1.0   5) gpu/0.17.3b  (g)
  2) slurm/expanse/21.08.8   4) DefaultModules   6) matlab/2022b/nmbr5dd
[mthomas@exp-9-55 ~]$ matlab
MATLAB is selecting SOFTWARE OPENGL rendering.
```

# HPC Interactive Jobs: Running Matlab with GUI

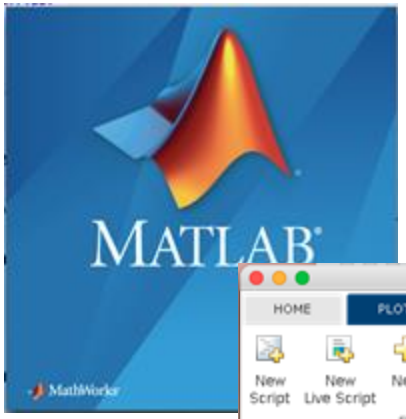


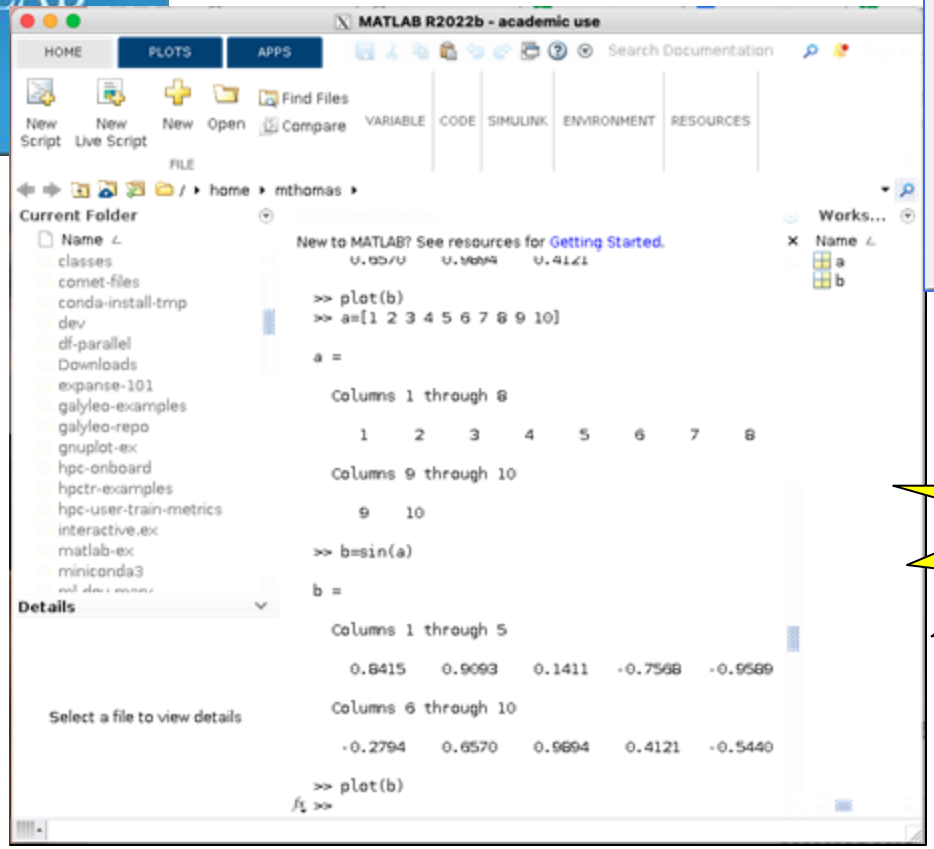


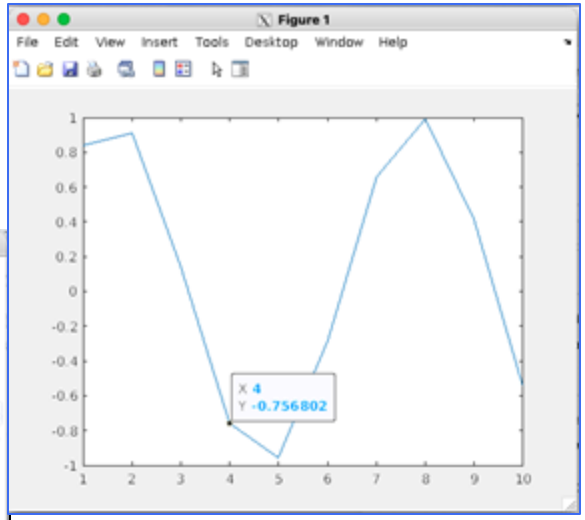


# Interactive Jobs: Running R (TSCC only)
# Expanse: use portal.expanse.sdsc.edu

***Step 3:*** *Request an interactive node, using* ***–X*** *option*

```
[mthomas@tscc-login2 ~]$ qsub -I -q glean -l nodes=1:ppn=1
qsub: waiting for job 26844488.tscc-mgr7.local to start
qsub: job 26844488.tscc-mgr7.local ready
[mthomas@tscc-4-46 ~]$
```

***Step 4:*** *Setup your module environment and run Matlab*


```
[mthomas@tscc-4-46 ~]$ module load R
Unloading compiler-dependent module openmpi_ib/3.1.4
[mthomas@tscc-4-46 ~]$ R
R version 4.0.2 (2020-06-22) -- "Taking Off Again"
Copyright (C) 2020 The R Foundation for Statistical Computing
Platform: x86_64-pc-linux-gnu (64-bit)

R is free software and comes with ABSOLUTELY NO WARRANTY.
[SNIP]
R is a collaborative project with many contributors.
[SNIP]
> myString <- "Hello, World!"
> print ( myString )
[1] "Hello, World!"
>
```

# Juptyer Notebooks

What is Jupyter?

> ***Jupyter*** *is a free, open-source, interactive web tool known as a computational notebook, which researchers can use to combine software code, computational output, explanatory text and multimedia resources in a single document.* *(J. Perkel, https://www.nature.com/articles/d41586-018-07196-1)*

Common Jupyter Services:

- Jupyter Notebooks (single user)
- JupyterLab: advanced version of notebook
- JupyterHub: multiuser Jupyter service

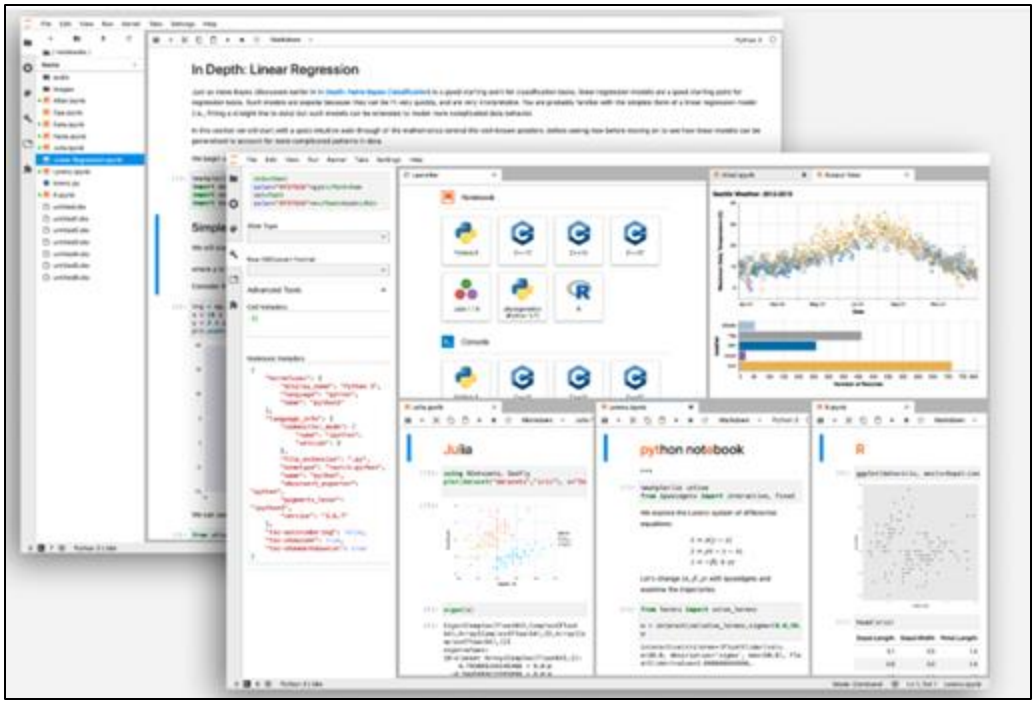

# Accessing and Running Secure Notebooks on SDSC HPC Systems

- Install notebook application:
  - Locally: install Anaconda on your laptop
  - Remotely:
    - Install Anaconda/conda on the remote machine (default is HTTP) – not recommended
- Running remotely:
  - Connect over HTTP (default, insecure)
  - Connect over HTTP + SSH tunneling (secure, but inconvenient)
  - Connect over HTTPS + using the *Satellite Reverse Proxy Service* (SRPS) and *galyleo* client (secure, convenient)
- ➢ You can launch Jupyter services on SDSC:
  - Launch securely (HTTPS) using SRPS/*galyleo* -- *recommended*
  - CPU and GPUs
  - Interactive nodes: command line *or* Slurm batch script
- Treat the Notebook URL like a Password!

# SDSC Satellite Reverse Proxy Service

- SRPS: prototype system that allows users to launch secure standard Jupyter Notebooks on on any Expanse compute node using a reverse proxy server.
  - Notebooks will be hosted on the internal cluster network as an HTTP service using standard Jupyter commands.
  - Service available to the user outside of the cluster firewall over HTTPS connection between the external users web browser and the reverse proxy server.
- Goal: minimize software changes for users, improve security of user notebooks running on SDSC systems.
- SRPS can run on any HPC system capable of supporting Apache on internal network.

# galyleo

- 2nd generation shell utility developed to orchestrate a user's interaction with both Satellite and Slurm to start a Jupyter session within a batch job.
- Developed while reviewing start-jupyter (prototype client) codebase to sort out how best to support Expanse (OOD) Portal and HPC User Services Group long-term; integrated into an existing SSH tunneling orchestration utility to use Satellite proxy service instead

- Key features in design:
  - HTTPS URL
  - Supports containers (Singularity on Expanse)  **GPU environment**
  - No need to install conda environment or update packages
  - Increases flexibility for users to configure software environment; but also try to makes it simpler for them to do this themselves
  - Batch job script is generated completely on-the-fly.
  - Command-line argument driven.
  - Quiet mode for OOD portal

https://github.com/mkandes/galyleo

# Running Notebooks Remotely & Securely Using SRPS

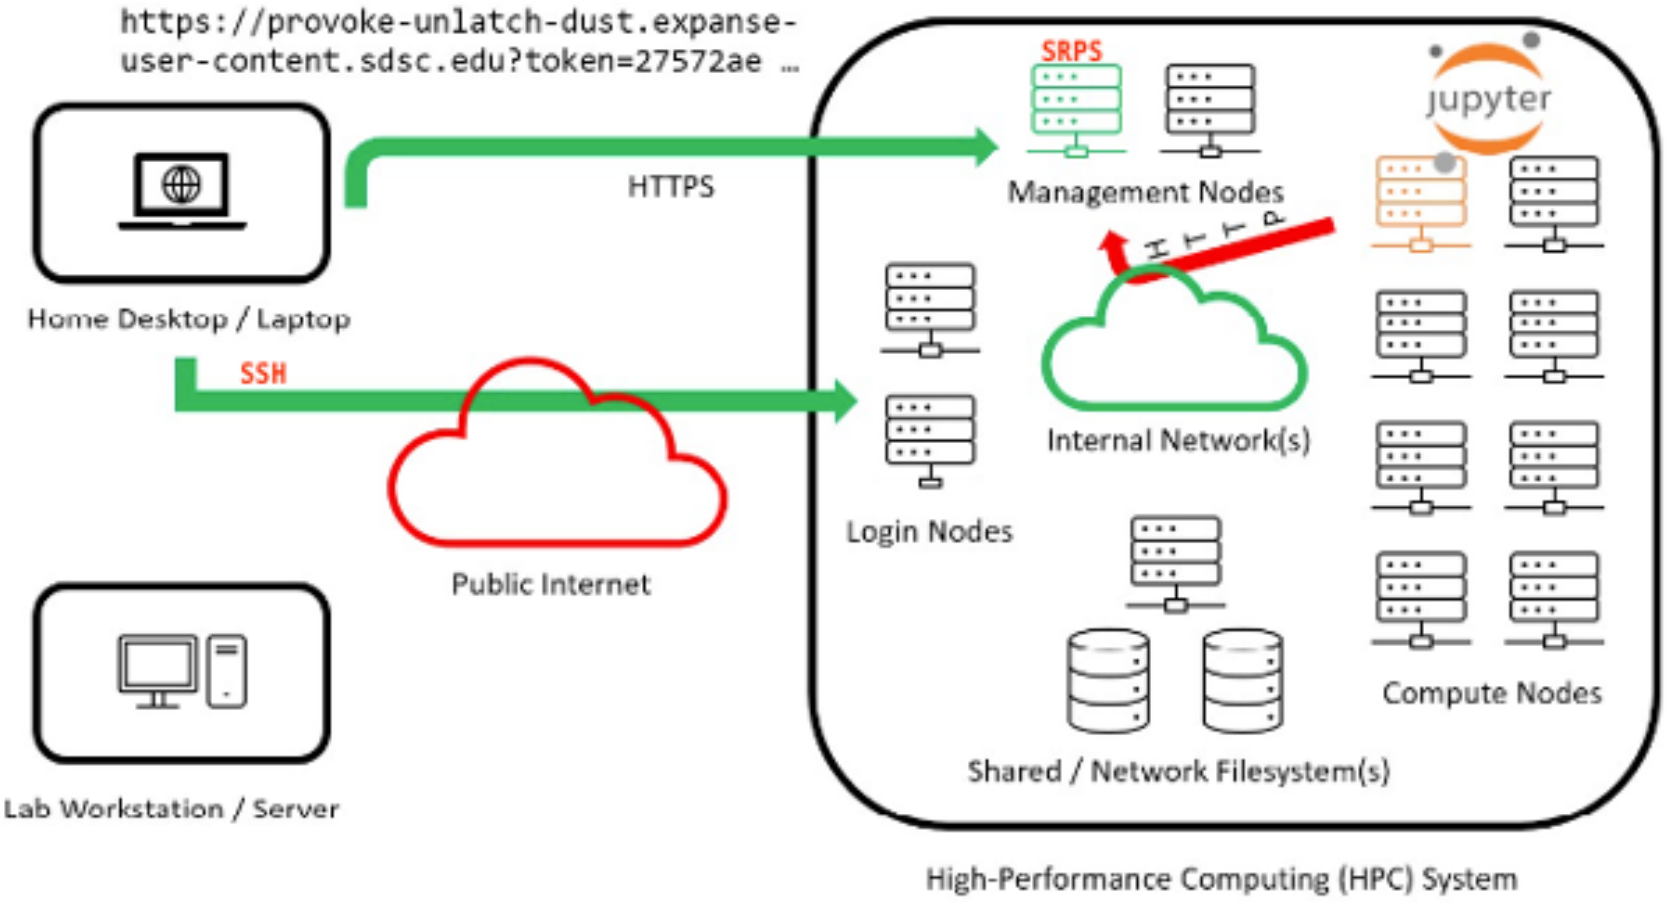


Src: M Kandes: https://education.sdsc.edu/training/interactive/202112_running_jupyter_notebooks_on_expanse/

# Satellite-Galyleo System

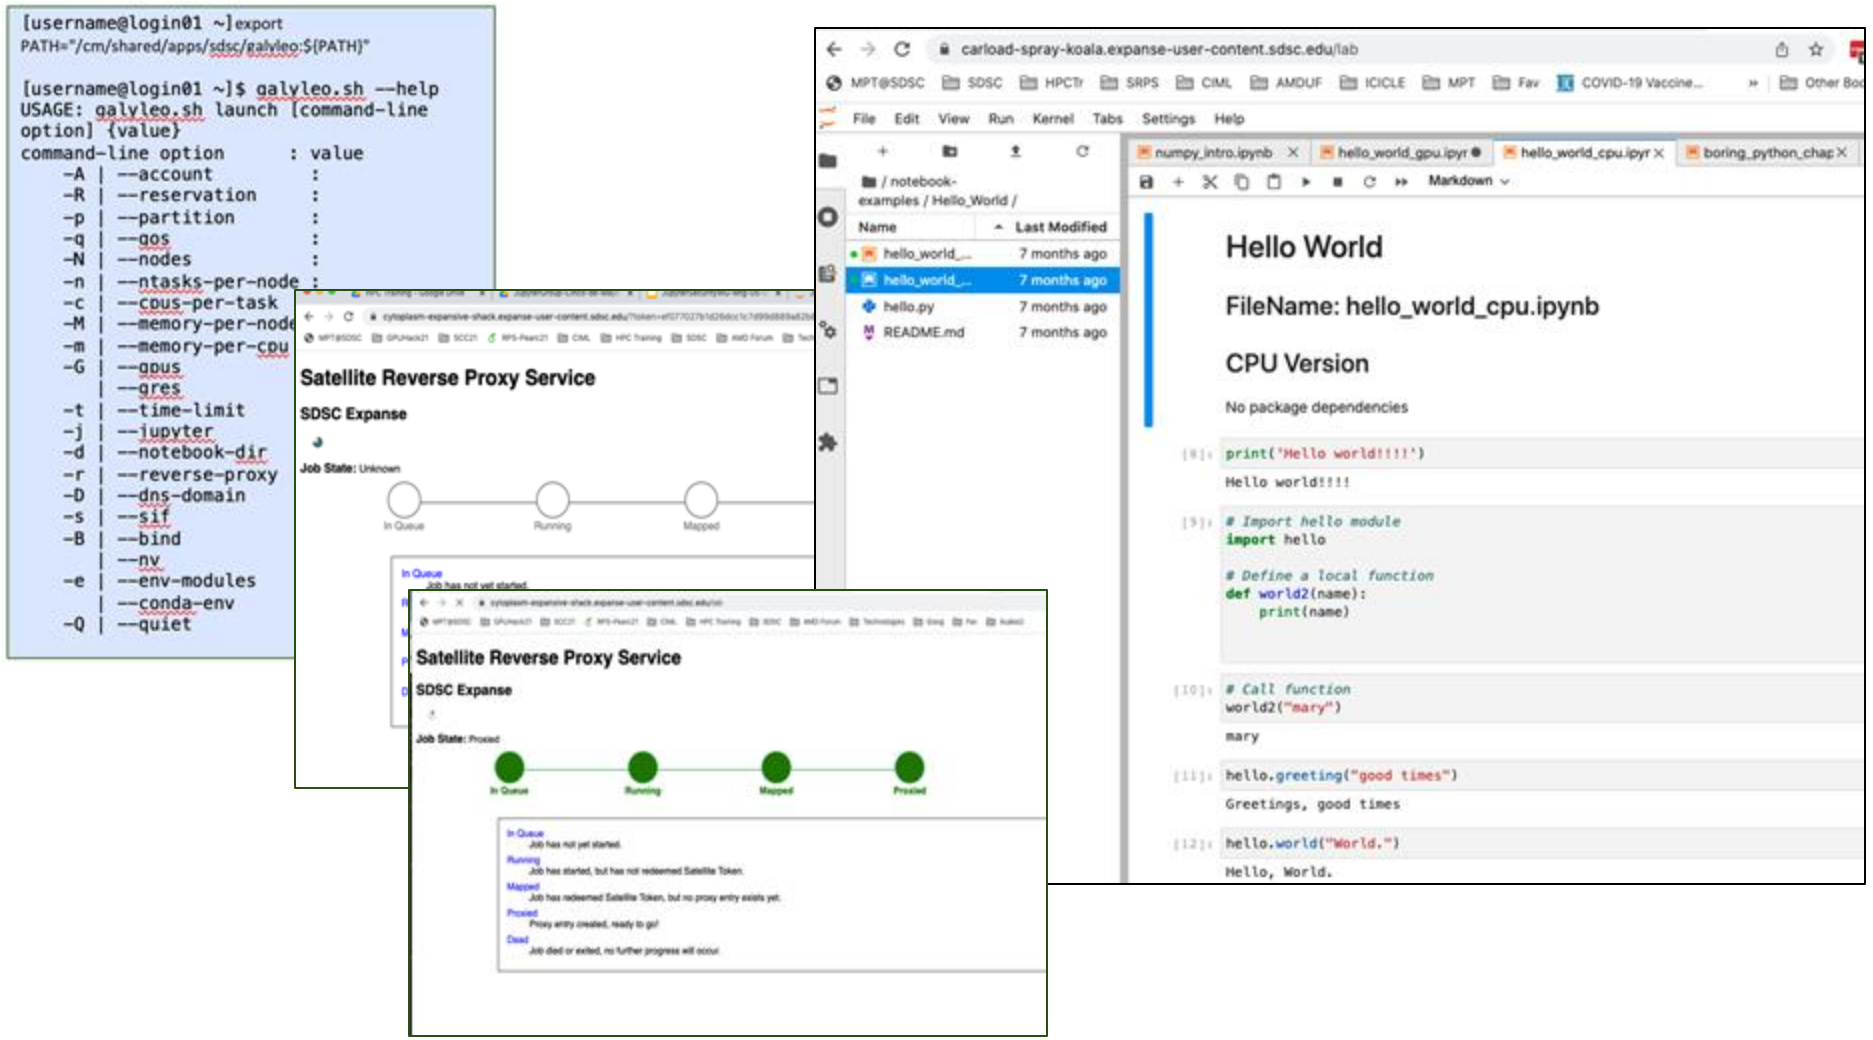

# Satellite Client: galyleo

Key features in design:

- User calls galyleo.sh launch script, which requests token from Satellite, passes token to batch job script and submits the job to Slurm; token redeemed from batch job once it runs
- Increase flexibility for users to configure software environment; but also try to make it simpler for them to do themselves
- Batch job script is generated completely on-the-fly.
- Command-line argument driven.
- Quiet mode for OOD portal

```
[username@login01 ~]export PATH="/cm/shared/apps/sdsc/galyleo:${PATH}"

[username@login01 ~]$ galyleo.sh --help
USAGE: galyleo.sh launch [command-line option] {value}
command-line option      : value
   -A | --account        :
   -R | --reservation    :
   -p | --partition      :
   -q | --qos            :
   -N | --nodes          :
   -n | --ntasks-per-node :
   -c | --cpus-per-task  :
   -M | --memory-per-node :  GB
   -m | --memory-per-cpu  :  GB
   -G | --gpus           :
      | --gres           :
   -t | --time-limit     :
   -j | --jupyter        :
   -d | --notebook-dir   :
   -r | --reverse-proxy  :
   -D | --dns-domain     :
   -s | --sif            :
   -B | --bind           :
      | --nv             :
   -e | --env-modules    :
      | --conda-env      :
   -Q | --quiet          : 1
```

https://github.com/mkandes/galyleo

# Launching CPU notebooks using galyleo

**Step 1:**
Login and setup user environment

**Step 2:**
Run command to launch secure notebook

**Step 3:**
Copy URL; paste into browser on local system

**Step 4:**
Monitor Slurm queue; wait for job to start

```
[username@login01 ~]$ which galyleo.sh
/usr/bin/which: no galyleo.sh in (/cm/shared/apps/spack/cpu/opt/spack/linux-centos8-zen/gcc-8.3.1
[SNIP]
home/username/.local/bin:/home/username/bin)
[username@login01 ~]$ export PATH="/cm/shared/apps/sdsc/galyleo:${PATH}"
[username@login01 ~]$ which galyleo
/cm/shared/apps/sdsc/galyleo/galyleo
[username@login01 ~]$ galyleo launch --account abc123 --partition shared --cpus 2 --memory 4 --time-limit
00:30:00 --env-modules cpu/0.17.3b,anaconda3/2021.05
[snip]
Submitted Jupyter launch script to Slurm. Your SLURM_JOB_ID is 9773665
[snip]
Your Jupyter notebook session will begin once compute resources are allocated to your job by the scheduler.

https://carload-spray-koala.expanse-user-content.sdsc.edu/lab

[username@login01 ~]$ squeue -u username
             JOBID PARTITION     NAME     USER ST       TIME  NODES NODELIST(REASON)
           9773665 gpu-debug galyleo-  username PD       0:00      1 (None)
[username@login01 ~]$ squeue -u username
             JOBID PARTITION     NAME     USER ST       TIME  NODES NODELIST(REASON)
           9773665 gpu-debug galyleo-  username  R       0:20      1 exp-7-59
```

# Satellite Server Status/Pending Page

- Load notebook URL in browser; wait for it to launch
- Monitor pending page
- Run the “squeue” command on the HPC system to check job status
- If the job queue is busy, it may take a while to launch the notebook
- Treat Jupyter Notebook URL as a password!

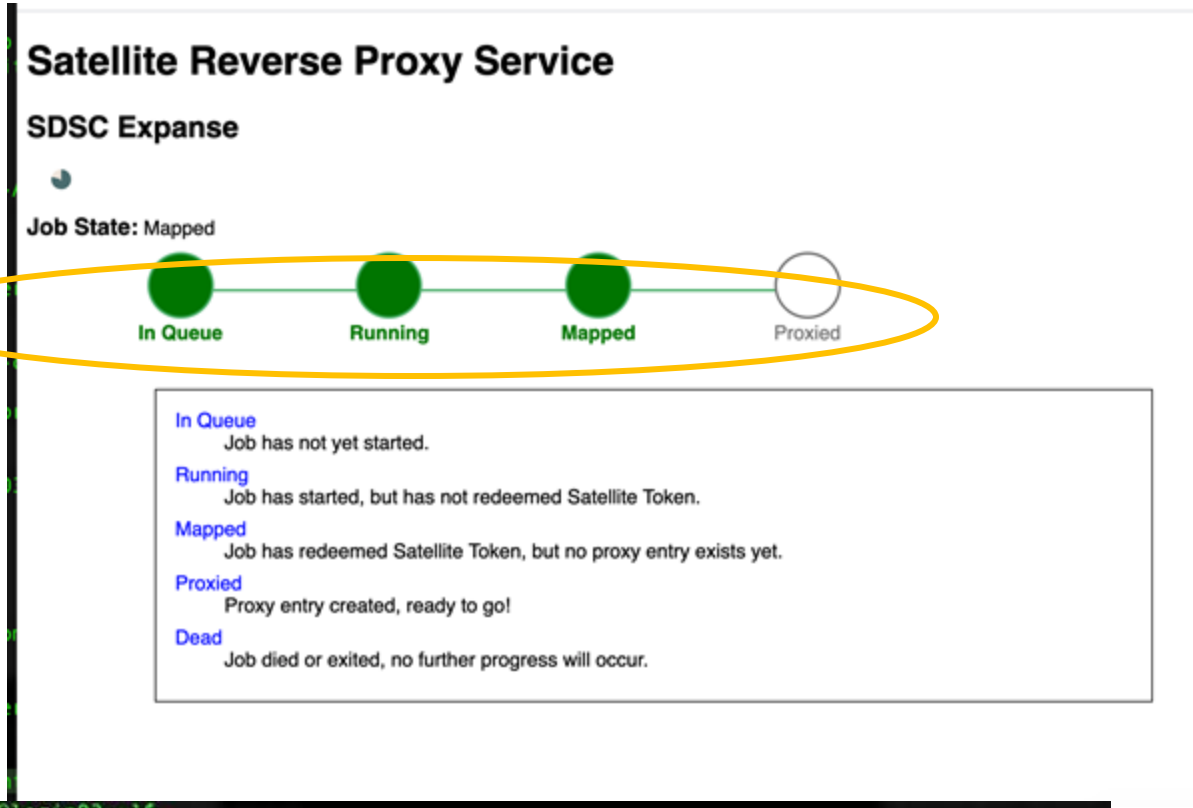


```
[mthomas@login02 ~]$
[mthomas@login02 ~]$ squeue -u mthomas
             JOBID PARTITION     NAME     USER ST       TIME  NODES NODELIST(REASON)
           9774239    shared galyleo-  mthomas  R       3:49      1 exp-1-13
           9774274     debug galyleo-  mthomas  R       0:12      1 exp-9-55
[mthomas@login02 ~]$
```

# Launching GPU notebooks using galyleo

- GPU Notebooks run better when using containers. SDSC maintains several containers on Expanse
- See: /cm/shared/apps/containers/singularity

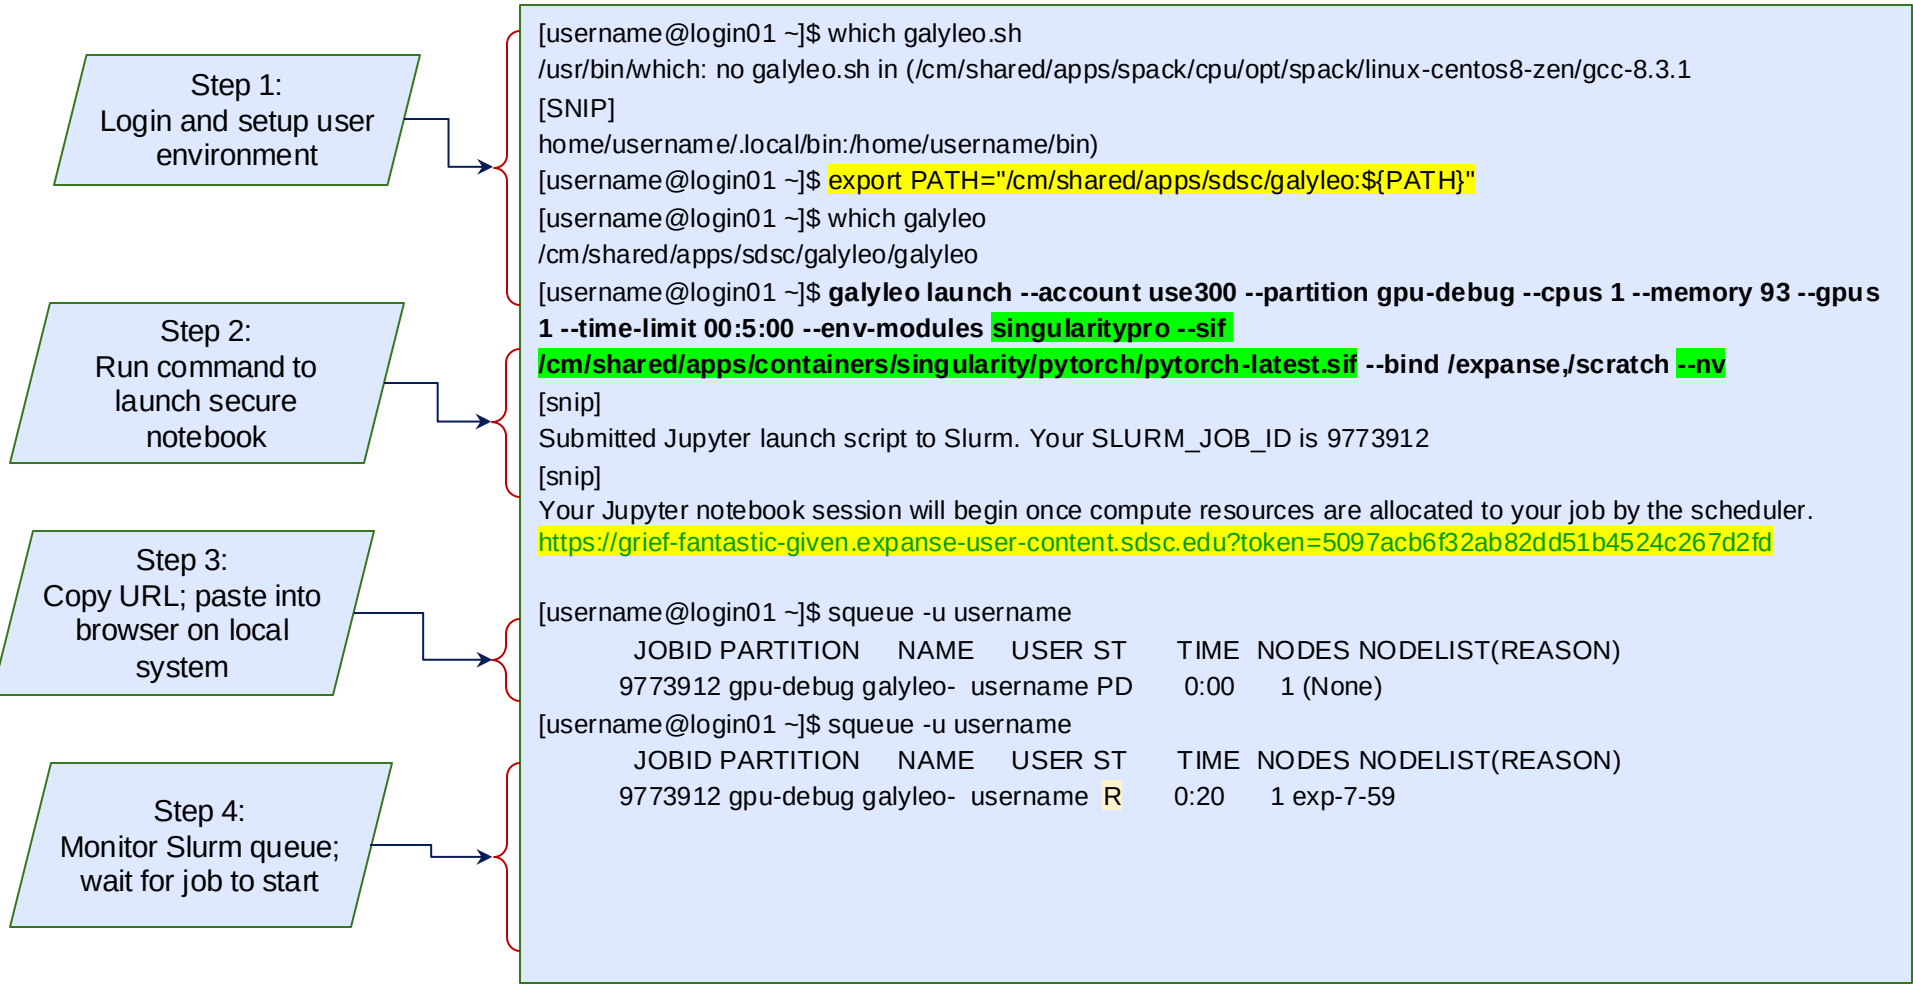

# Verify notebook launched on GPU device

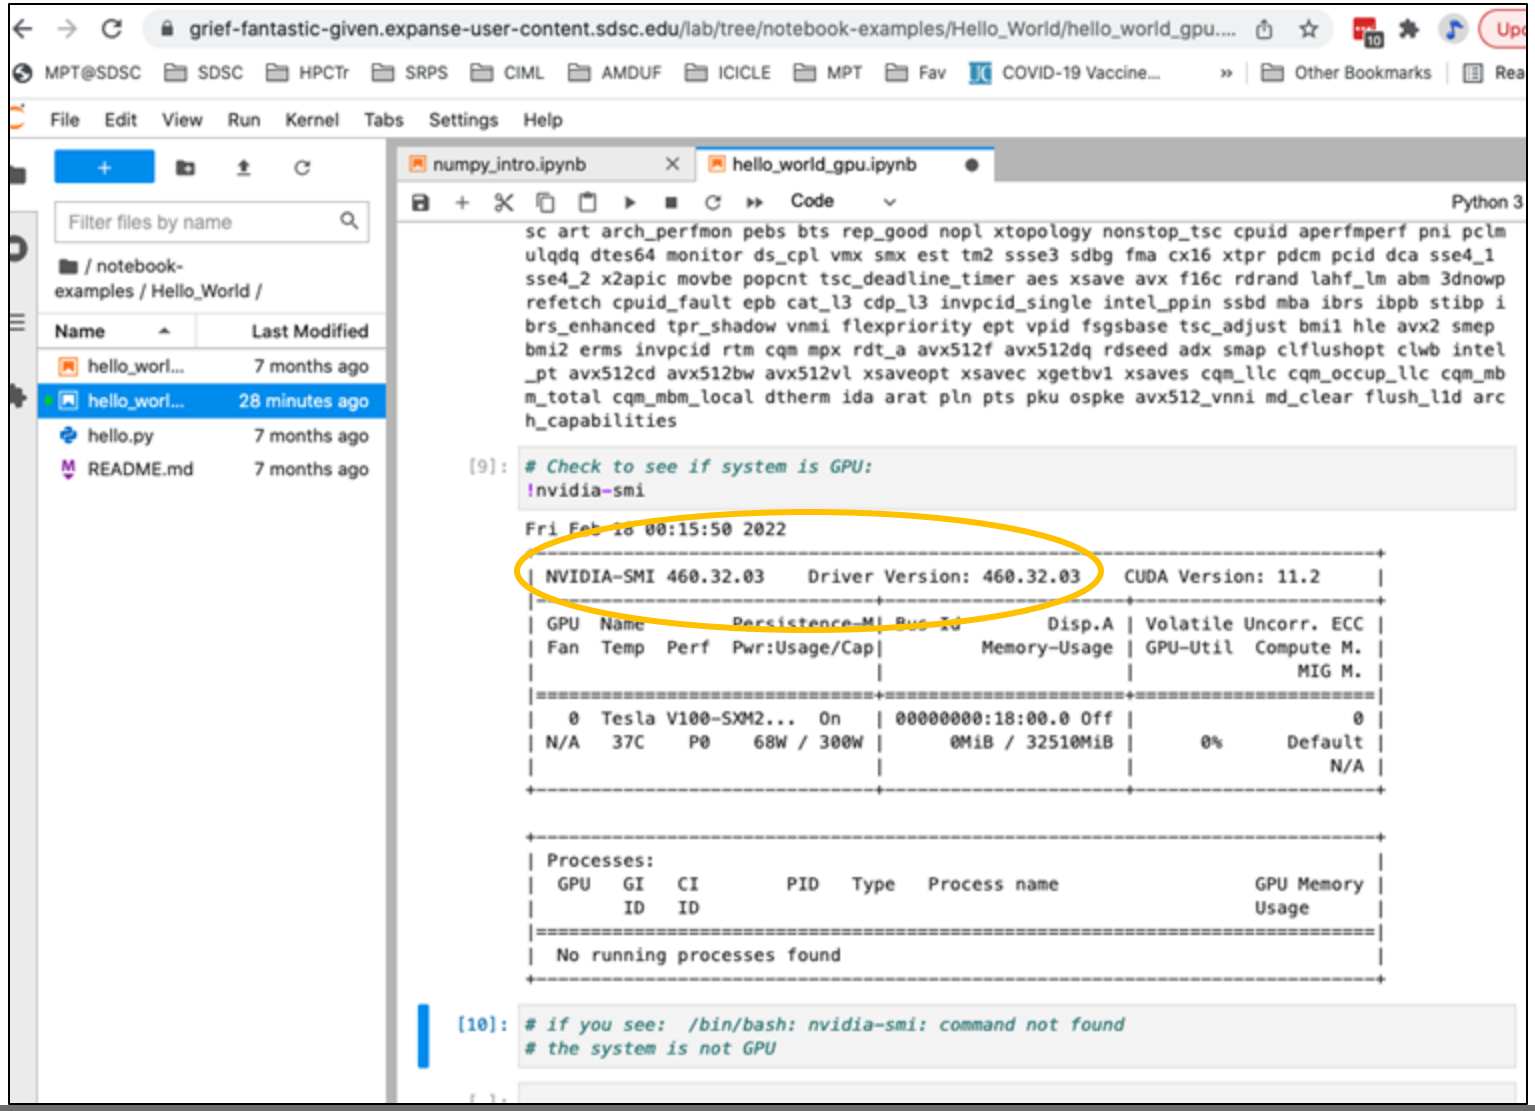

# Notebook Examples

https://github.com/sdsc-hpc-training-org/Expanse-Notebooks

- Collection of notebooks tested on Expanse and other SDSC HPC systems
- Includes range of materials from “hello world” to ML notebooks.
- View by name, CPU/GPU, or Serial/Parallel

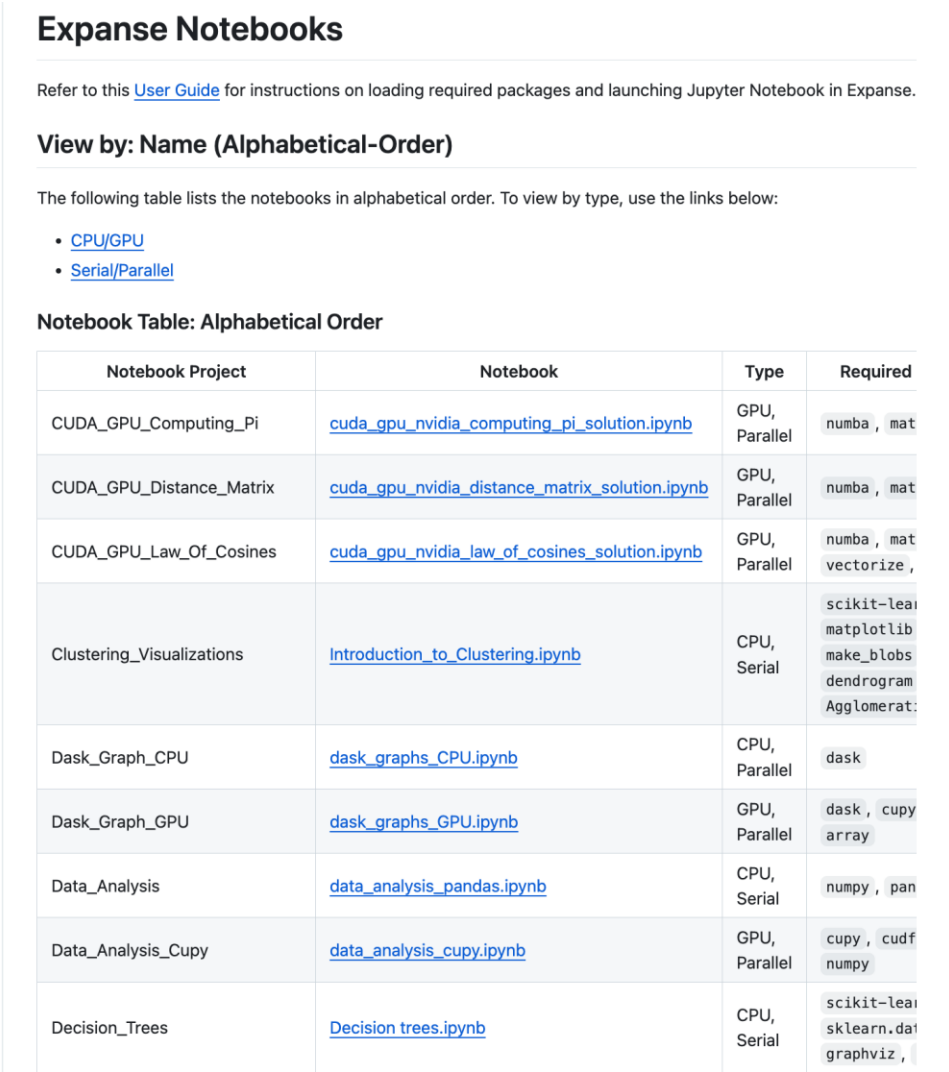

## Expanse Notebooks

Refer to this User Guide for instructions on loading required packages and launching Jupyter Notebook in Expanse.

### View by: Name (Alphabetical-Order)

The following table lists the notebooks in alphabetical order. To view by type, use the links below:

- CPU/GPU
- Serial/Parallel

### Notebook Table: Alphabetical Order

| Notebook Project | Notebook | Type | Required |
|---|---|---|---|
| CUDA_GPU_Computing_Pi | cuda_gpu_nvidia_computing_pi_solution.ipynb | GPU, Parallel | numba, mat |
| CUDA_GPU_Distance_Matrix | cuda_gpu_nvidia_distance_matrix_solution.ipynb | GPU, Parallel | numba, mat |
| CUDA_GPU_Law_Of_Cosines | cuda_gpu_nvidia_law_of_cosines_solution.ipynb | GPU, Parallel | numba, mat vectorize, |
| Clustering_Visualizations | Introduction_to_Clustering.ipynb | CPU, Serial | scikit-lea matplotlib make_blobs dendrogram Agglomerat |
| Dask_Graph_CPU | dask_graphs_CPU.ipynb | CPU, Parallel | dask |
| Dask_Graph_GPU | dask_graphs_GPU.ipynb | GPU, Parallel | dask, cupy array |
| Data_Analysis | data_analysis_pandas.ipynb | CPU, Serial | numpy, pan |
| Data_Analysis_Cupy | data_analysis_cupy.ipynb | GPU, Parallel | cupy, cudf numpy |
| Decision_Trees | Decision trees.ipynb | CPU, Serial | scikit-lea sklearn.dat graphviz, |

# Outline

- Defining Interactive HPC
  - High-performance computing (HPC)
  - HPC batch computing
  - Interactive computing
- Accessing Interactive HPC Nodes
  - Launching nodes
  - Running GUIs using X11 forwarding
- Interactive Application Examples
  - Viewing Data: unix file ops (grep, awk, cat), gnuplot, NetCDF
  - Programming & Visualization Platforms: Matlab, R, Jupyter Notebooks
  - Gateways & Portals: simplify access to interactive apps
- Q&A

# Expanse User Portal

SIMPLIFY! USE THE EXPANSE PORTAL

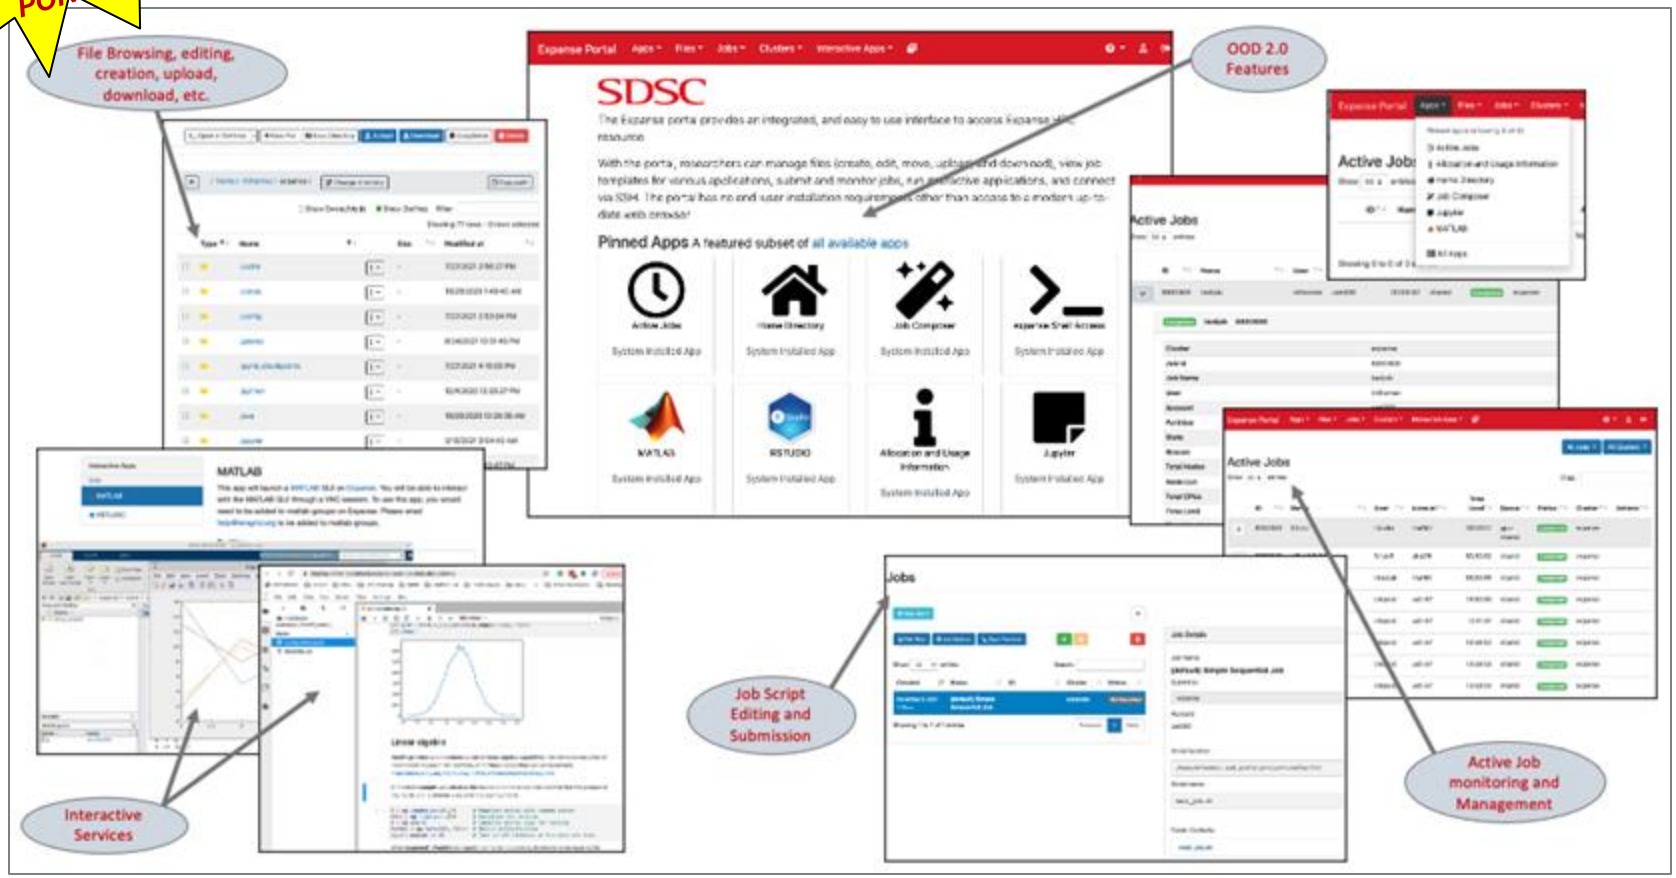


- https://portal.expanse.sdsc.edu; authenticate using ACCESS credentials
- Securely hosts batch job submission & monitoring, and interactive applications
- Simplifies launching supported interactive applications  manages software dependencies




Sivagnanam, S., Irving, C., Kandes, M., Mishin, D., Sakai, S., Strande, S., Tatineni, M., Thomas, M., Norman, M., (2021). Experiences in building a user portal for Expanse supercomputer. *PEARC21: Practice and Experience in Advanced Research Computing*, Updated: ACM. https://doi.org/10.1145/3437359.3465590

COMPLECS NSF award #2320934

# Expanse Portal: File Management

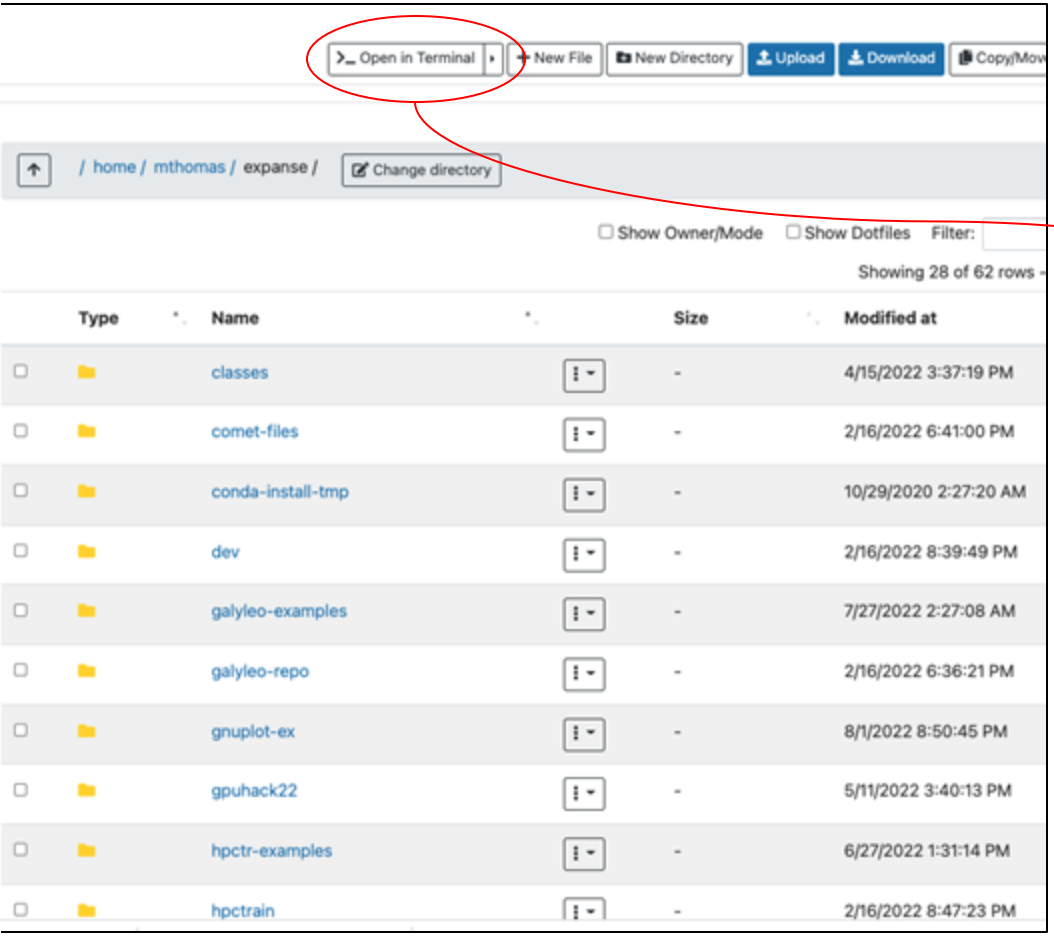


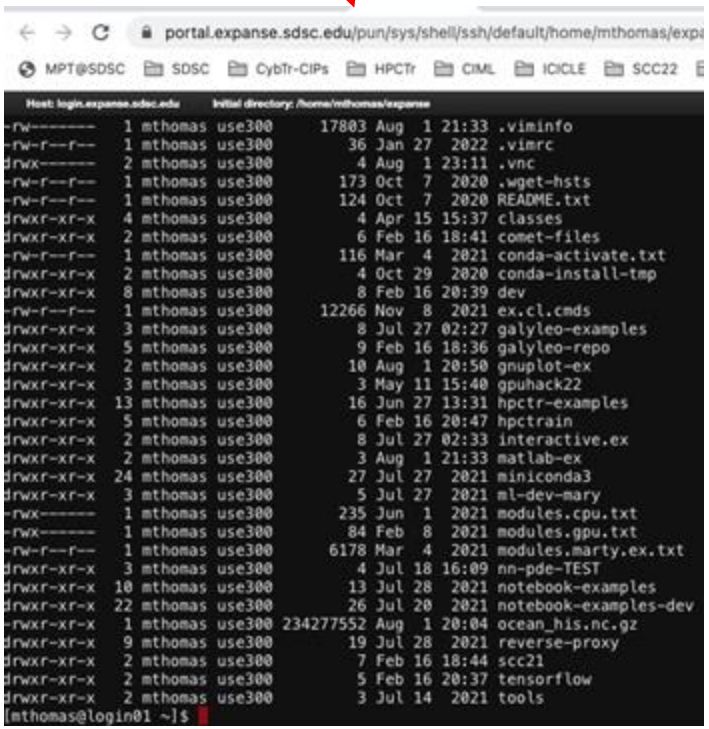

# Expanse Portal: Running Matlab

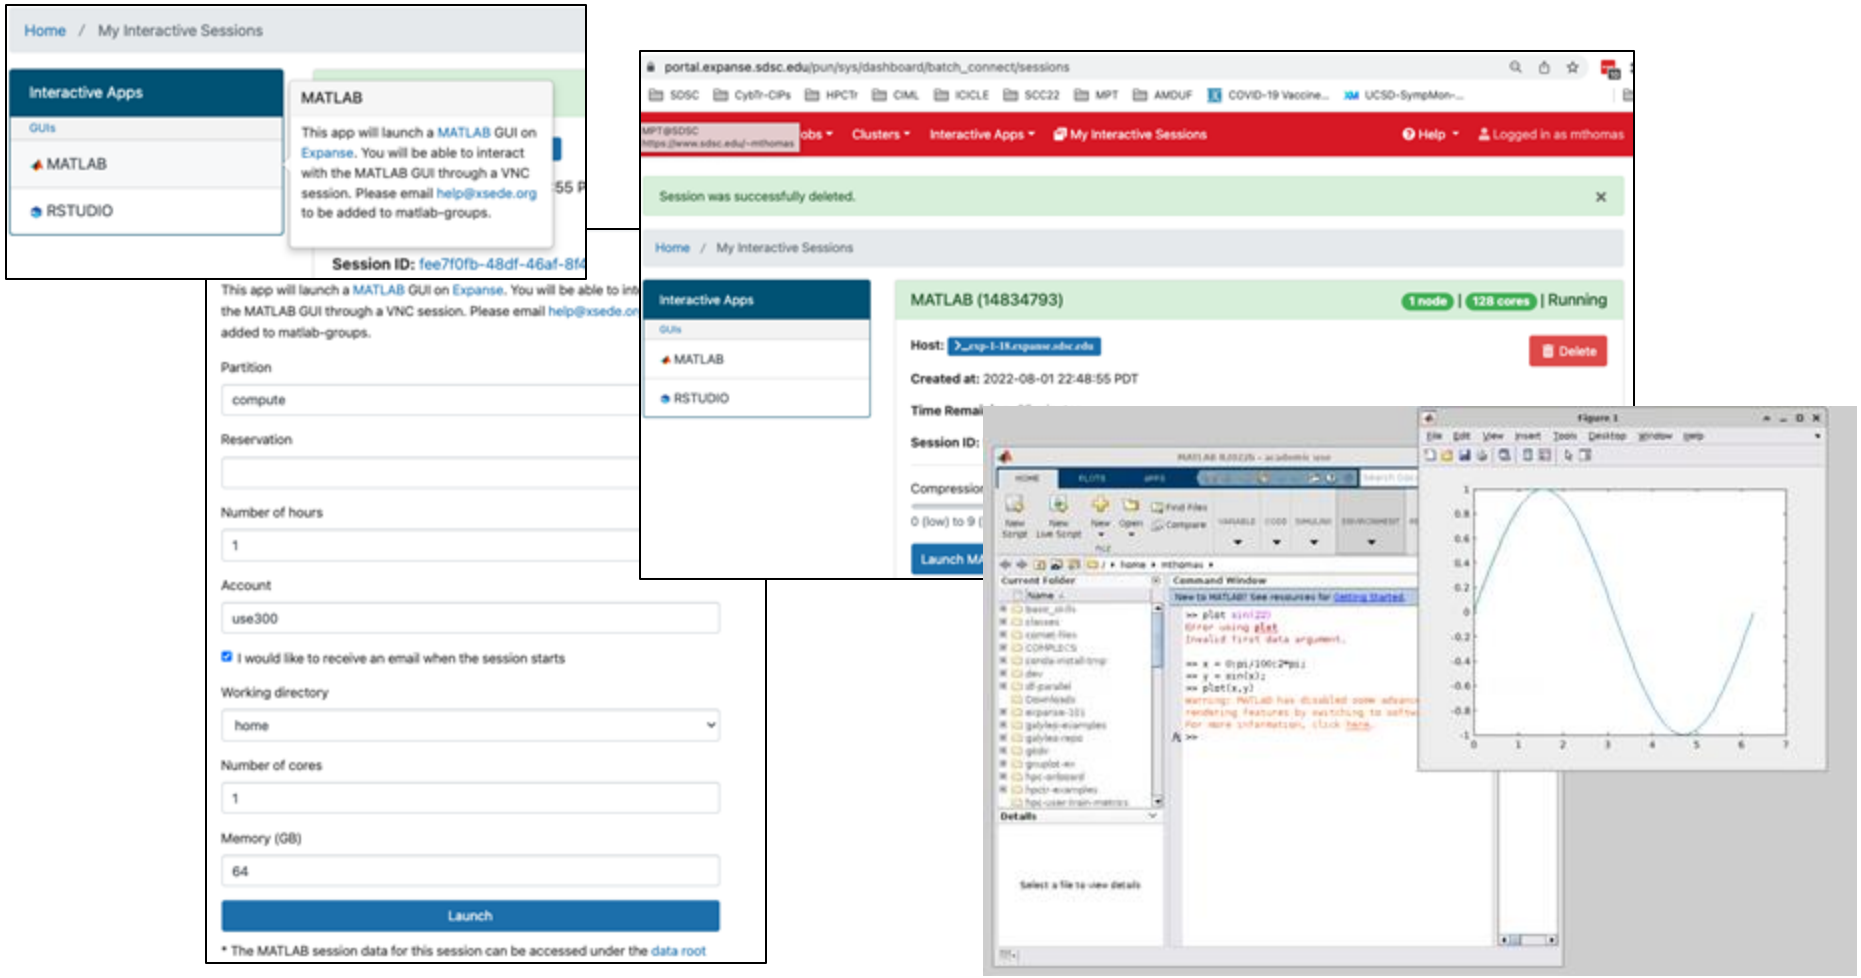

# Expanse Portal: Launching Notebooks

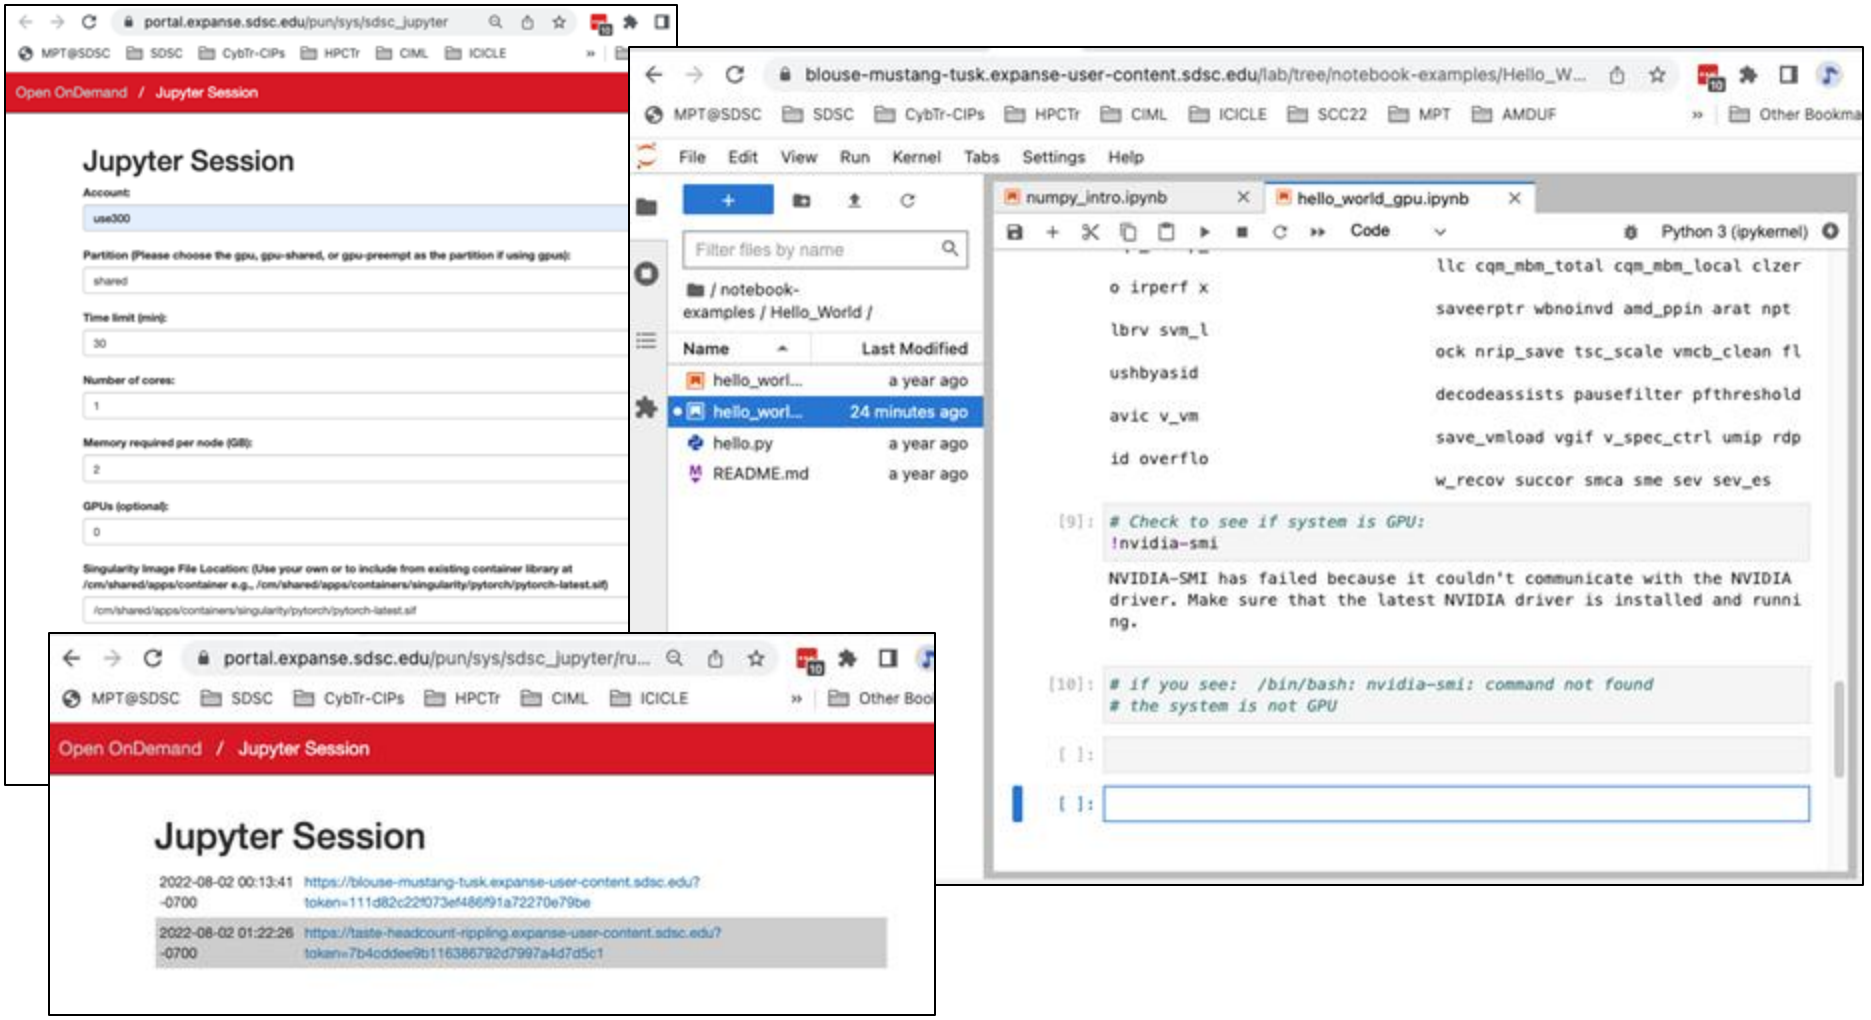

# Expanse Portal: Input data example for Jupyter Notebook

**Account:** use300
**Partition:** *(Please choose the gpu, gpu-shared, or gpu-preempt as the partition if using gpus): debug*
*Time limit (min)*: 30
**Number of cores:** 1
**Memory required per node (GB):** 2
**GPUs (optional):** 0
**Singularity Image File Location:** *(Use your own or to include from existing container library at /cm/shared/apps/container e.g., /cm/shared/apps/containers/singularity/pytorch/pytorch-latest.sif)*
/cm/shared/apps/containers/singularity/pytorch/pytorch-latest.sif
**Environment modules to be loaded** *(E.g., to use latest version of system Anaconda3 include cpu,gcc,anaconda3):* singularitypro
**Conda Environment** (Enter your own conda environment if any):
**Conda Init** (Provide path to conda initialization scripts):
**Conda Yaml** *(Upload a yaml file to build the conda environment at runtime) No file chosen:*
***Turn on use of mamba for speeding up conda-yml installs:***
***Enable use of new caching mechanism that will store and reuse conda-yml created environments using conda-pack !????***
**Reservation:**
**QoS:**
**Working directory:** HOME
**Type:** JupyterLab

# Outline

- Defining Interactive HPC
  - High-performance computing (HPC)
  - HPC batch computing
  - Interactive computing
- Accessing Interactive HPC Nodes
  - Launching nodes
  - Running GUIs using X11 forwarding
- Interactive Application Examples
  - Viewing Data: unix file ops (grep, awk, cat), gnuplot, NetCDF
  - Programming & Visualization Platforms: Matlab, R, Jupyter Notebooks
  - Gateways & Portals: simplify access to interactive apps
- Q&A

# Thank You!

# Q&A

If you have problems, please contact **consult@sdsc.edu**

https://github.com/sdsc-complecs/interactive-computing/

# Resources

- GitHub Repo for this presentation:
  - https://github.com/sdsc-complecs/interactive-computing/
- SDSC Training Resources
  - https://www.sdsc.edu/education_and_training/training
  - Code: https://github.com/sdsc-hpc-training-org/hpctr-examples
  - Running notebooks
    - using *galyleo:* https://github.com/mkandes/galyleo
    - Examples: https://github.com/sdsc-hpc-training-org/notebook-examples
- Expanse :
  - Landing page: expanse.sdsc.edu
  - User Guide: https://expanse.sdsc.edu/support/user_guides/expanse.html
  - Training: https://www.sdsc.edu/education_and_training/training_hpc.html
- Problems? Contact consult@sdsc.edu