


Deep Learning Topics: Special Connections Part I Encoder-Decoders

Paul Rodriguez, PhD
(SDSC)
June 2024



Outline

- **Part I**

Gate connection idea

Skip and Residual connections

Programing connections and Keras Model API

Encoder-Decoder (Autoencoder)

Exercise MNIST Autoencoder


Autoencoder with Stable Diffusion

- **Part II**

Attention Head and Transformers,



1. Gate connection: Recall the logistic unit

$$f(x) = \frac{1}{1 + \exp(-(b + wx))} \quad \longleftrightarrow \quad \begin{array}{c} b \\ \swarrow \\ x \xrightarrow{w} \bigcirc \rightarrow \text{output value} \end{array}$$

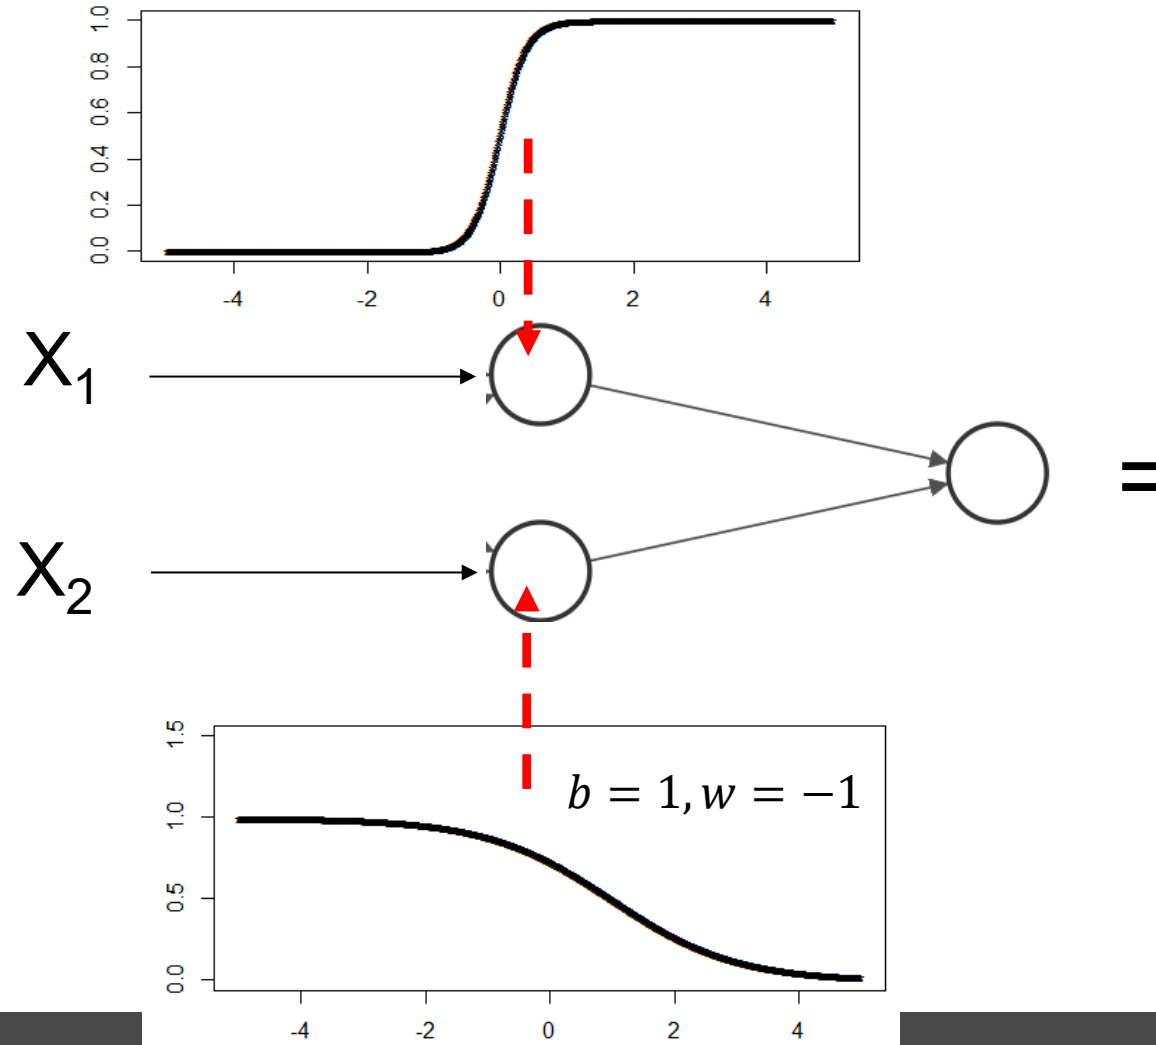


Example: 1 input into 2 logistic units with these activations

If you add these 2 units into a final output unit what would the output function look like?




Example: 2 input into 2 logistic units with these activations




Example: 2 input into 2 logistic units with these activations

*What if you multiply these?
What is the output function doing?*



Example: 2 input into 2 logistic units with these activations



What if you multiply these?


For linear activation, what is the output function doing ?

$$= \begin{cases} 0 & \text{if } X_1 < 0 \\ h(X_2) & \text{if } X_1 > 0 \end{cases}$$

X_1 "gates" X_2 activation

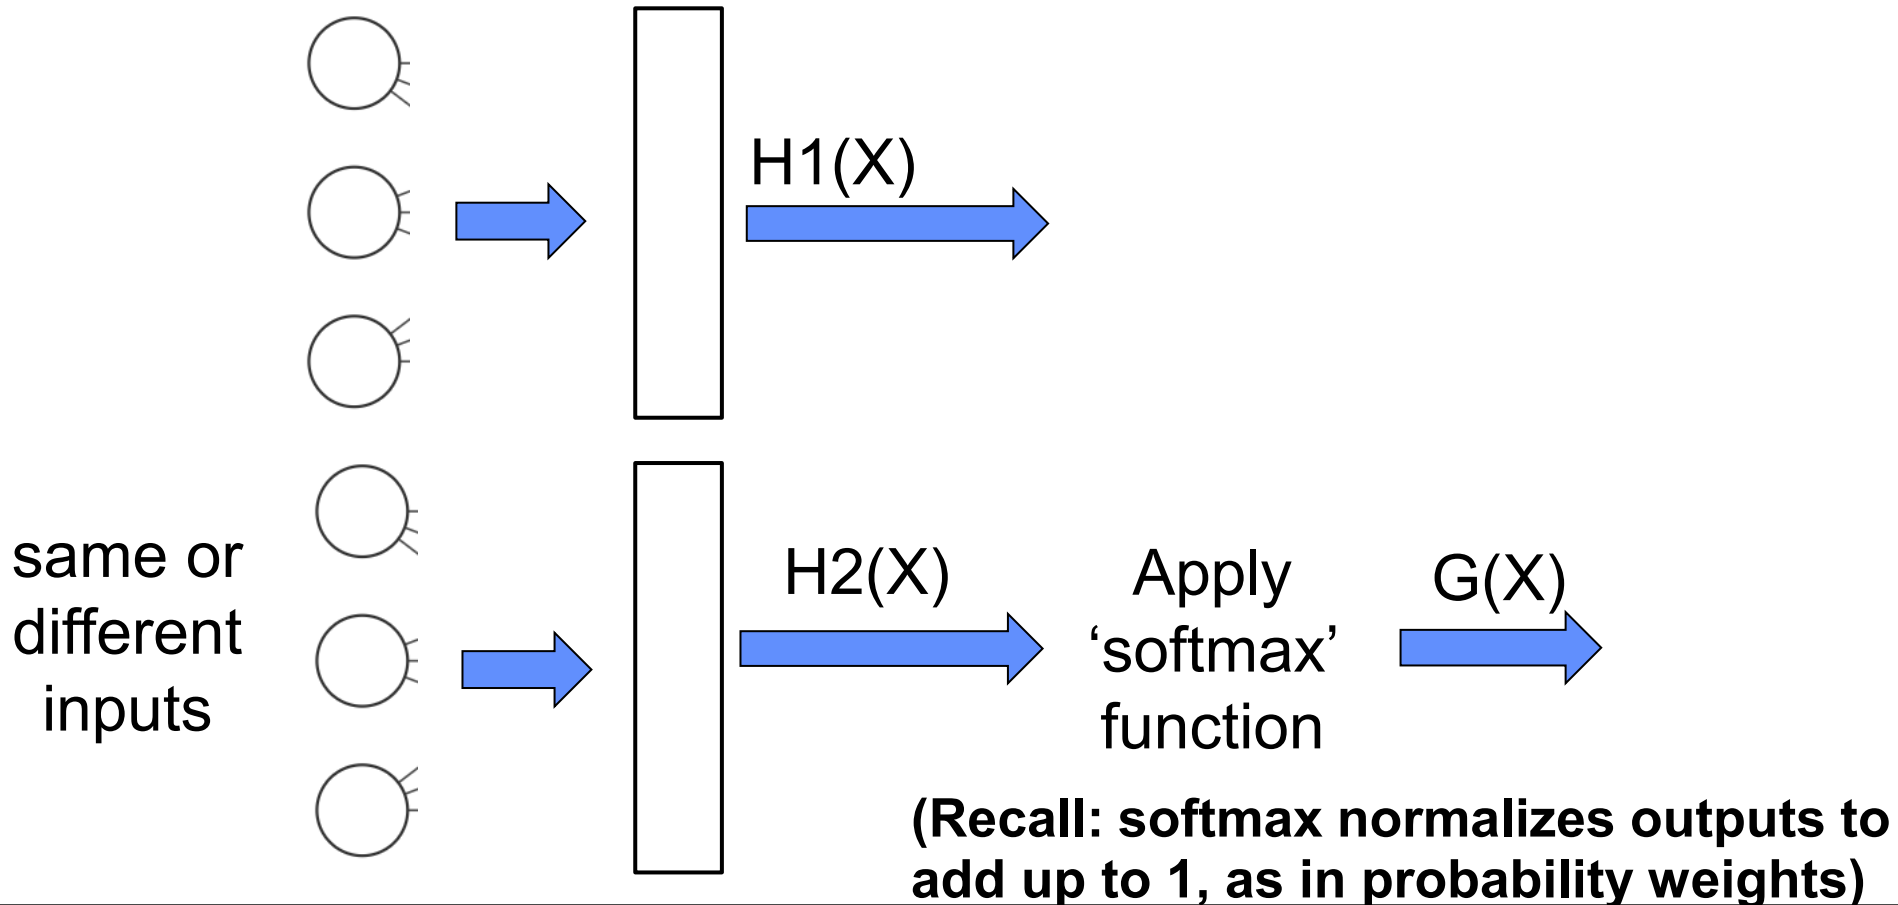
Drawing the gate for two sets of hidden units

Input layer Hidden layer



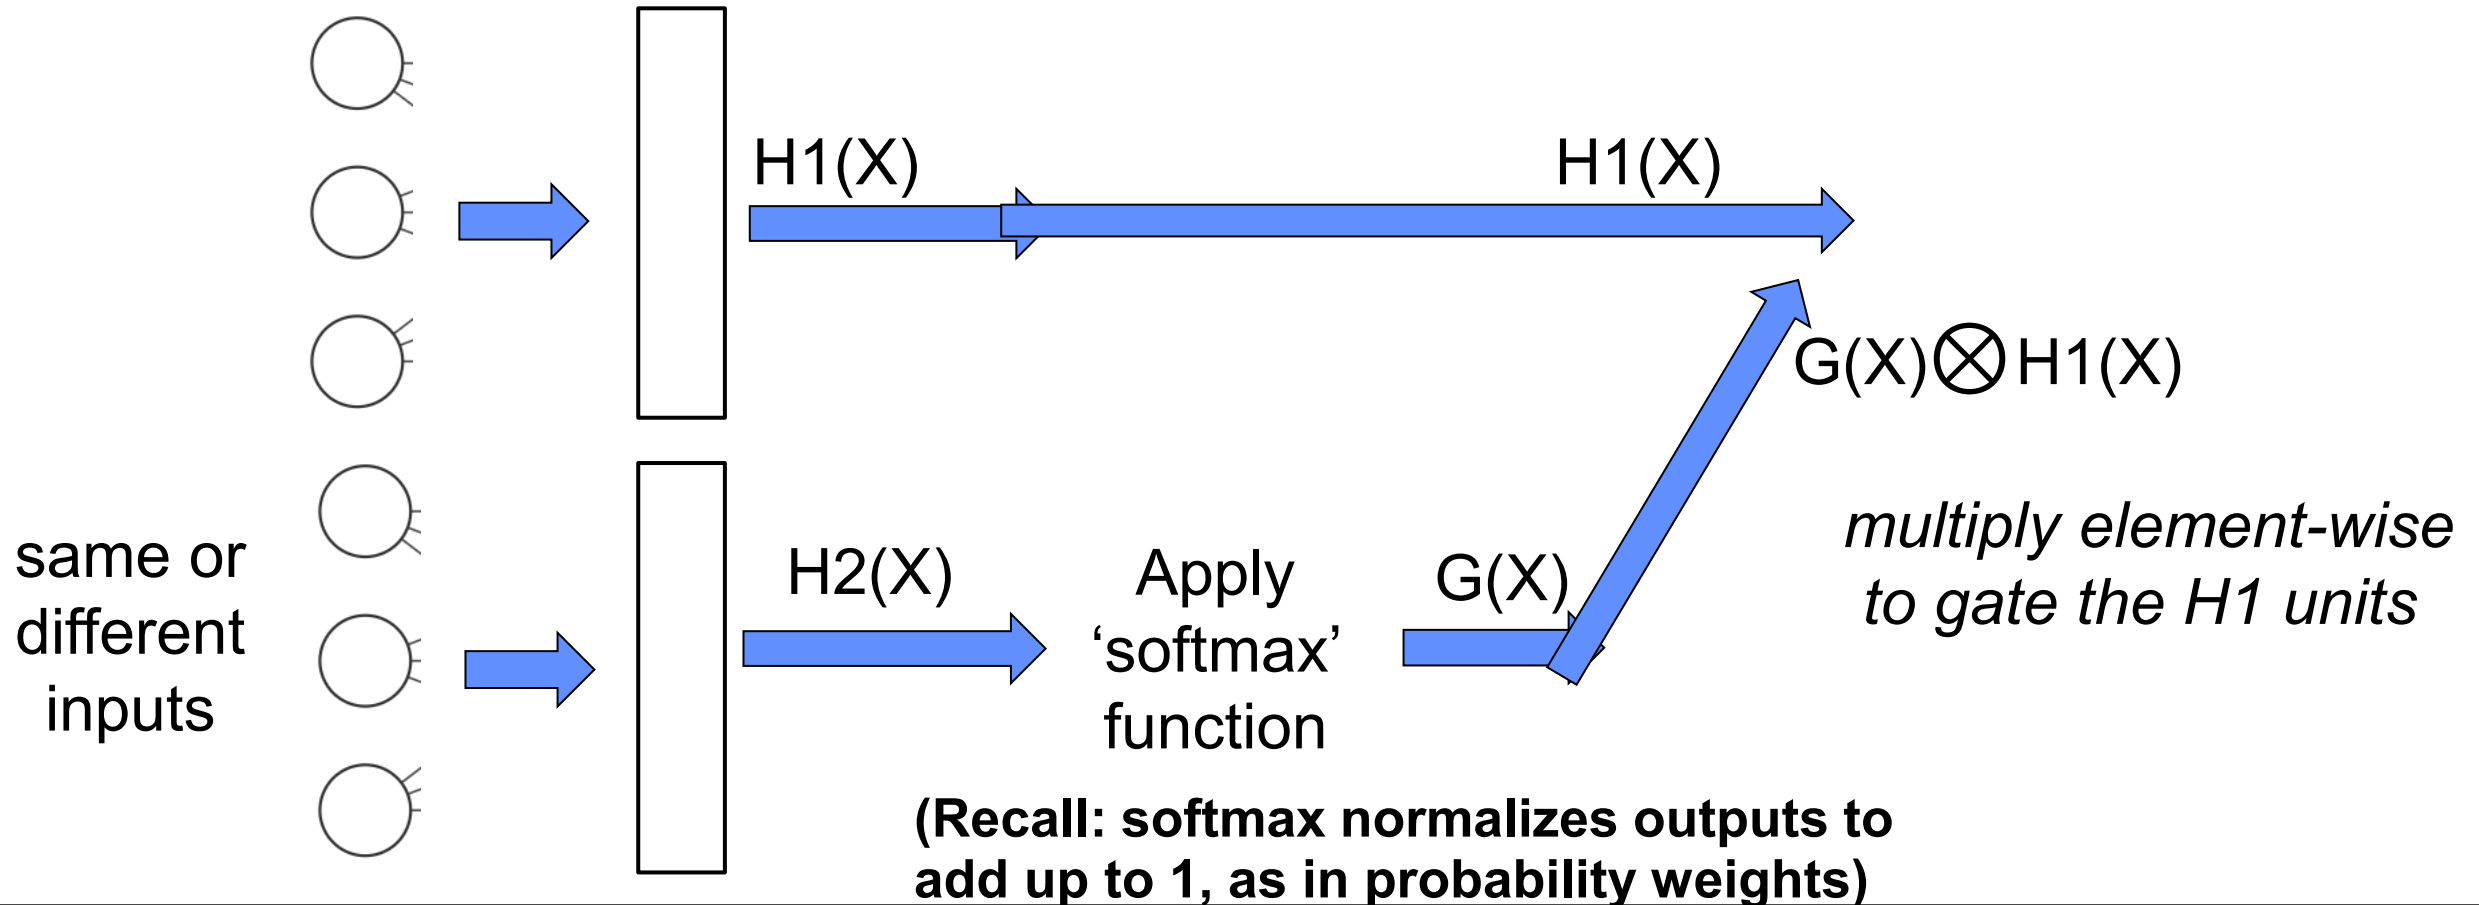
Drawing the gate for two sets of hidden units

Input layer Hidden layer




Drawing the gate for two sets of hidden units

Input layer Hidden layer



Scaled Dot-Product Attention (very rough summary)

“Attention” mechanism in language transformers use a softmax gate




The gate is applied to possible prediction Values for decoding

Q,K,V depend on input


Vaswani, et al. 2017
Attention Is All You Need (for Transformers)

2. Skip connection: Recall the Multilayer Perceptron




To help the MLP learn directly from input carry input forward

Input layer Hidden layer Output layer




Concatenate input with hidden units into new layer



*carry input forward
No weights to learn*

*Now the input
features are used
more efficiently*

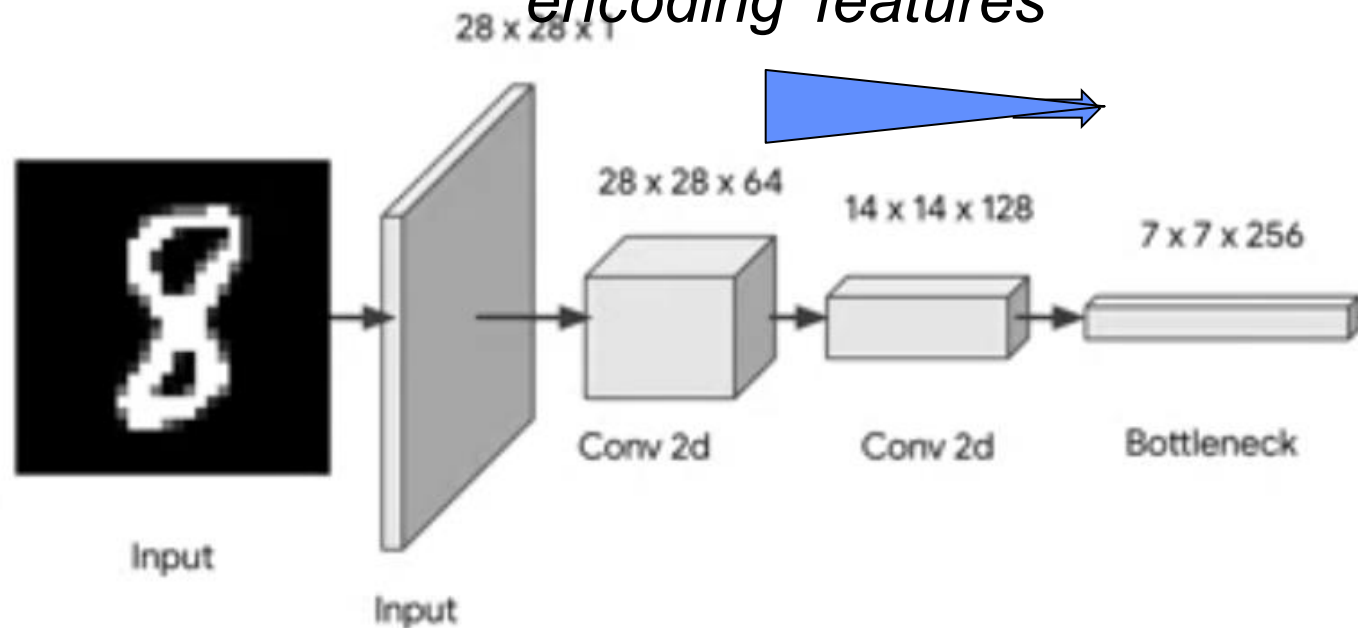
Can be done for any (or all) previous layer and *skip* any number of layers



Recall: CNN architecture for MNIST classification

ENCODER

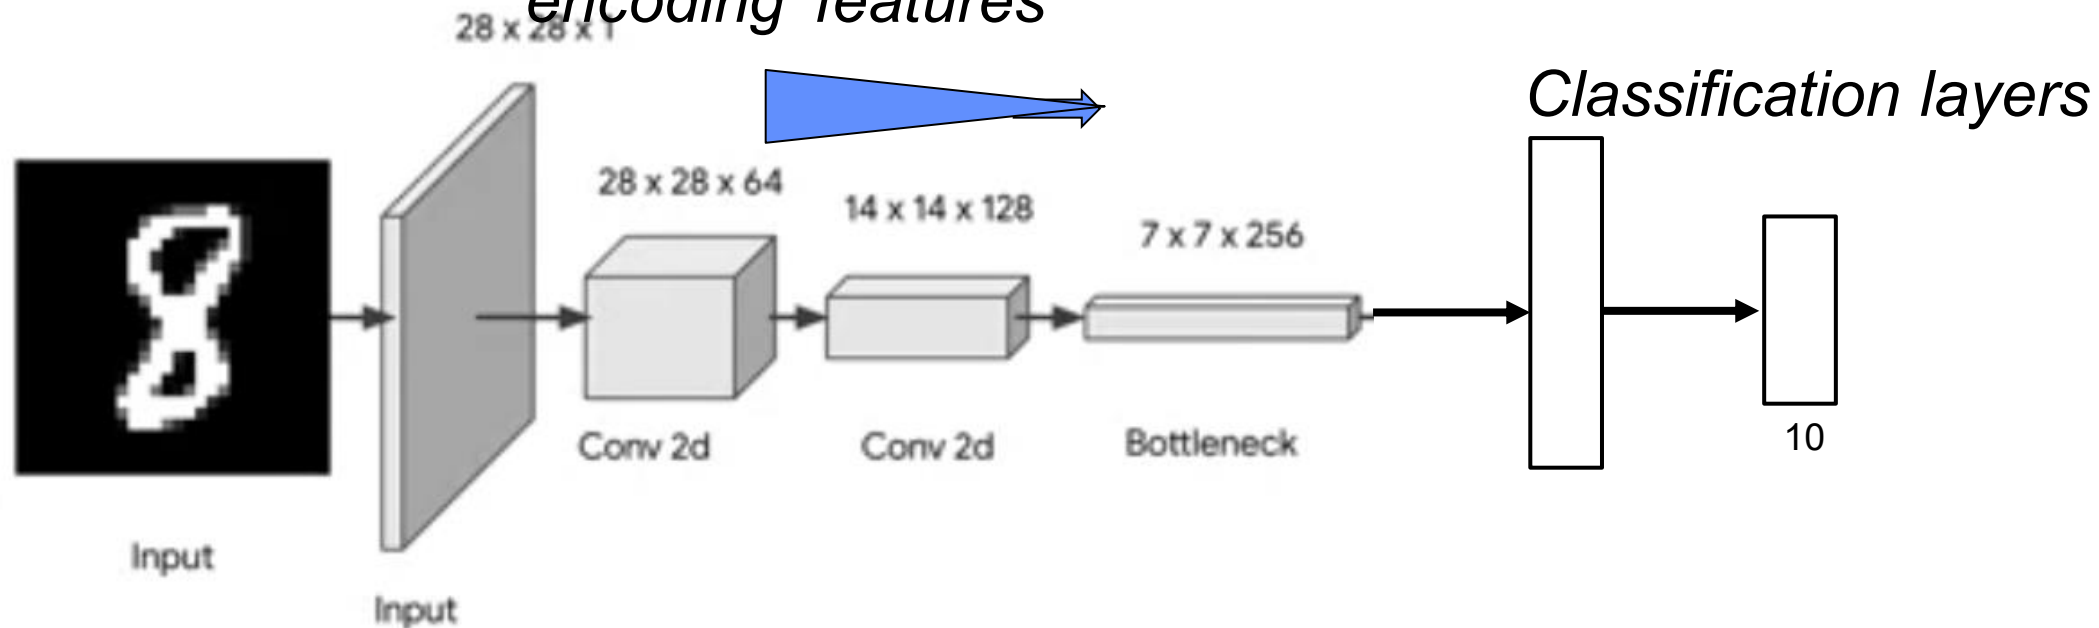
*more feature maps & downsampling :
'encoding' features*



Consider: CNN architecture for MNIST classification

ENCODER


*more feature maps & downsampling :
'encoding' features*



A CNN architecture for MNIST autoencoding

ENCODER

*more feature maps & downsampling :
'encoding' features*



*What if we want to
reconstruction the input?*


A CNN architecture for MNIST autoencoding

ENCODER

DECODER

*more feature maps & downsampling :
'encoding' features*

*What if we want to
reconstruction the input?*

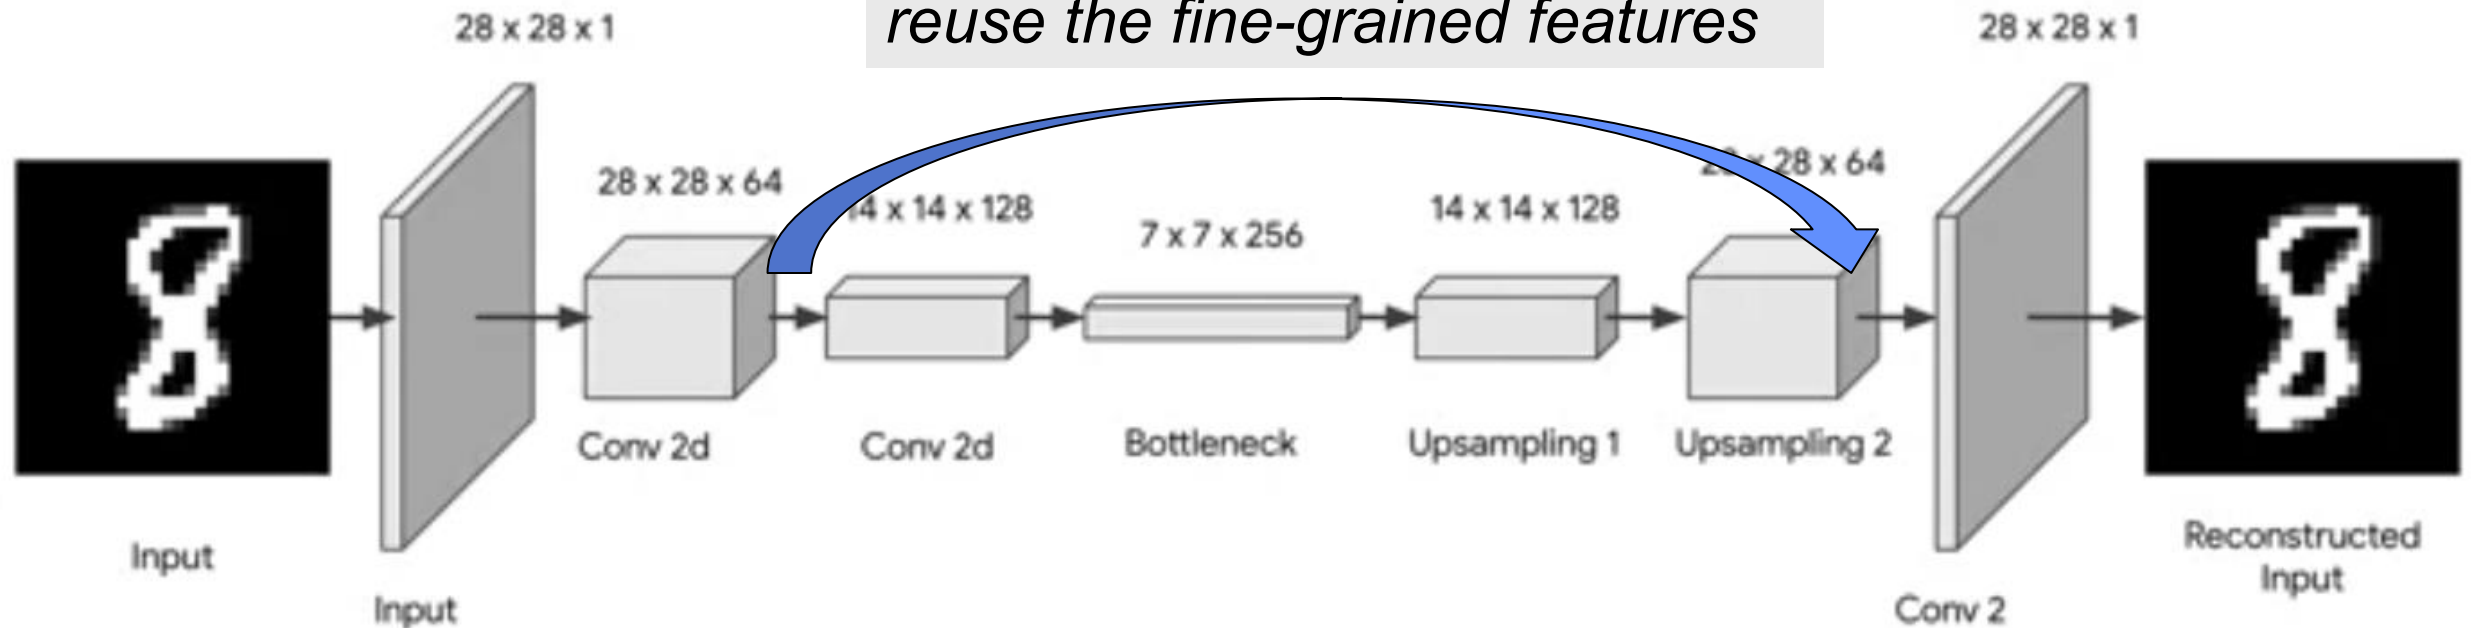


A CNN architecture for MNIST autoencoding

ENCODER

DECODER

Adding skip connections helps reuse the fine-grained features

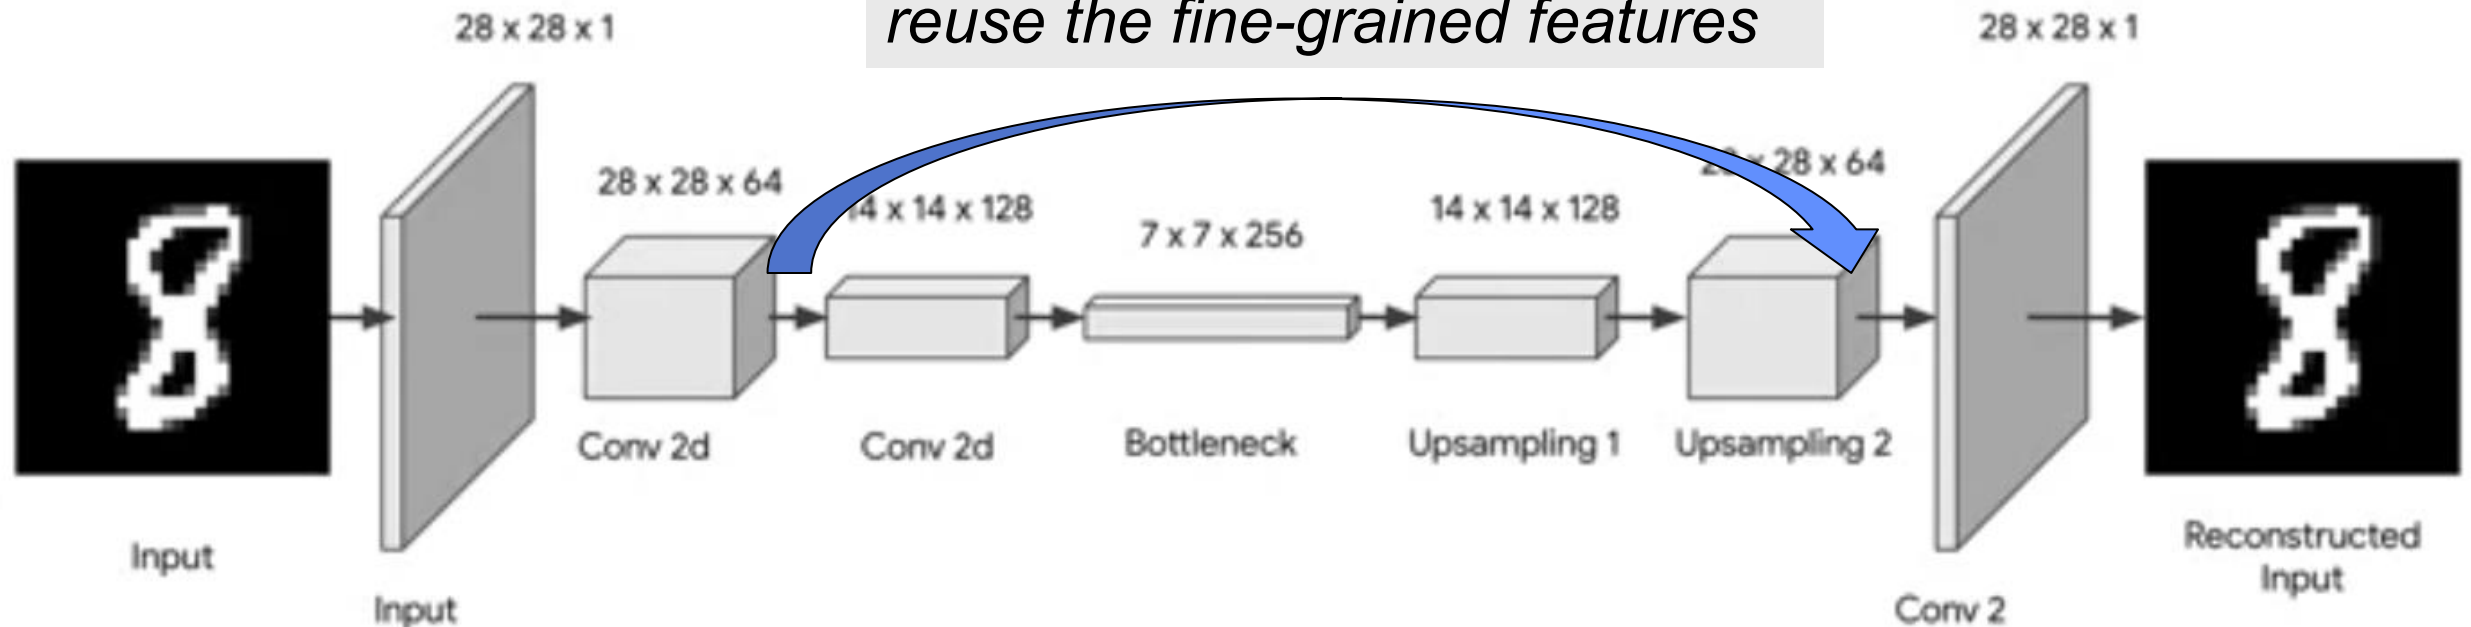


A CNN architecture for MNIST autoencoding

ENCODER


DECODER

Adding skip connections helps reuse the fine-grained features




NOTE the 28x28x64 encoded maps have to be skipped ahead to where the 28x28x64 decoding maps are – which axis is concatenated?

Image Encoder-Decoder is a “UNET” architecture




- **pause**

3. Residual connection: Can we keep adding layers?




Given some deep network,
should I add another layer?
What should a new layer learn?

Consider: Can we keep adding layers?



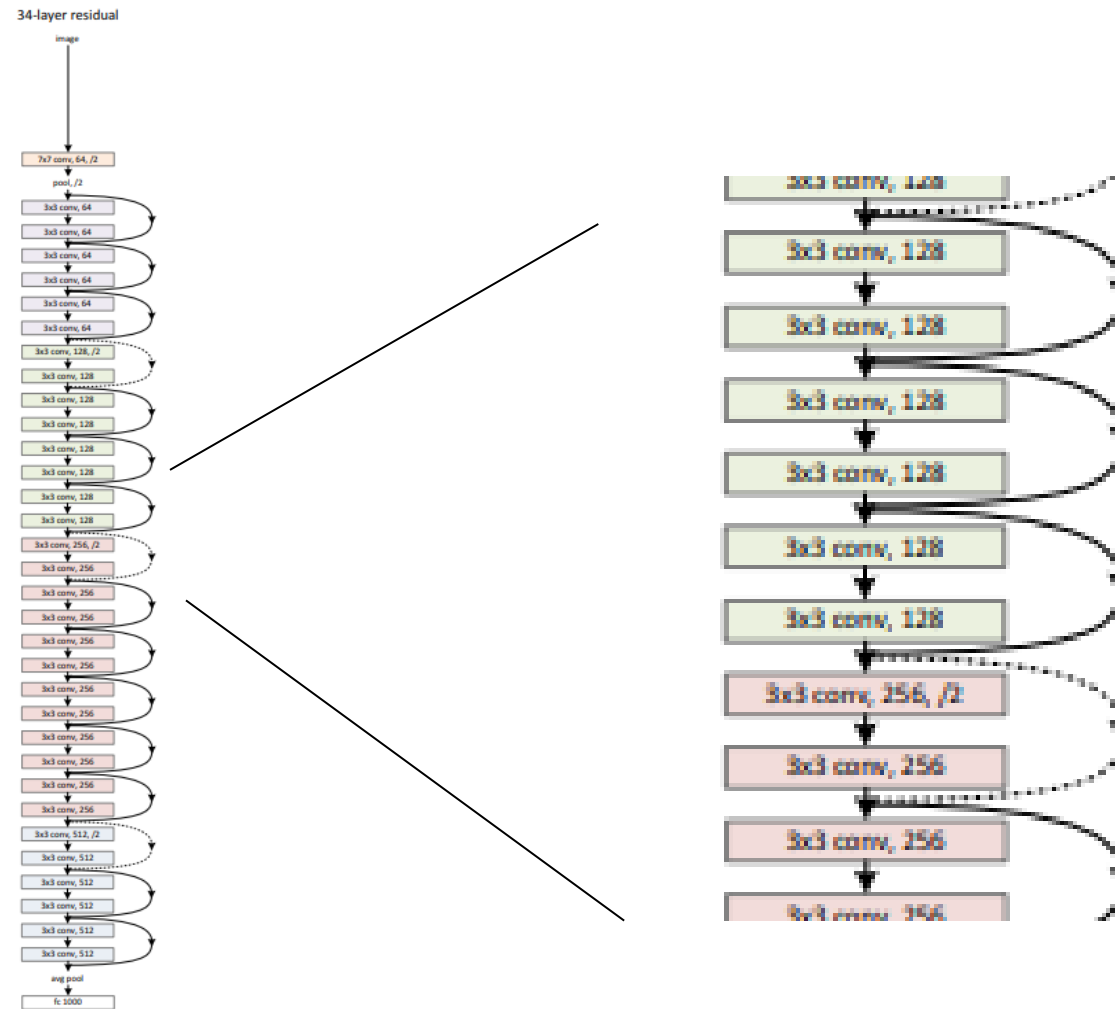
If $H1(X)$ is good then this new layer could be unnecessary or only for learning small differences
- eg $H2(X)$ should be almost same as $H1(X)$

Skip with addition makes a 'residual' connection



Make it easy for next layer to learn only incremental changes –
e.g. use $F(X)=H2(X)+H1(X)$ so that $H2(X)=F(X)-H1(X)$.
The $H2()$ function learned is a residual function


“Resnet” residual connections help deeper learning



*Deep Residual Learning,
He et.al, 2015*


Summary: useful connections for architectures, and the intuitions

Softmax for gating




Recurrent nets,
language transformer nets

Skip connections
for feature reuse



UNET, also
feedforward nets..

Residual connections
help deeper learning



Resnet, large image
classification

Exercise

- **MNIST autoencoder, reconstruct digits from noisy inputs**
- **Add skip connections with concatenation**

Note: make sure you see how the outputs from encoding layers are matched up to inputs for decoding layers!

e.g. 14x14 encoding feature maps should be concatenated with 14x14 decoding maps

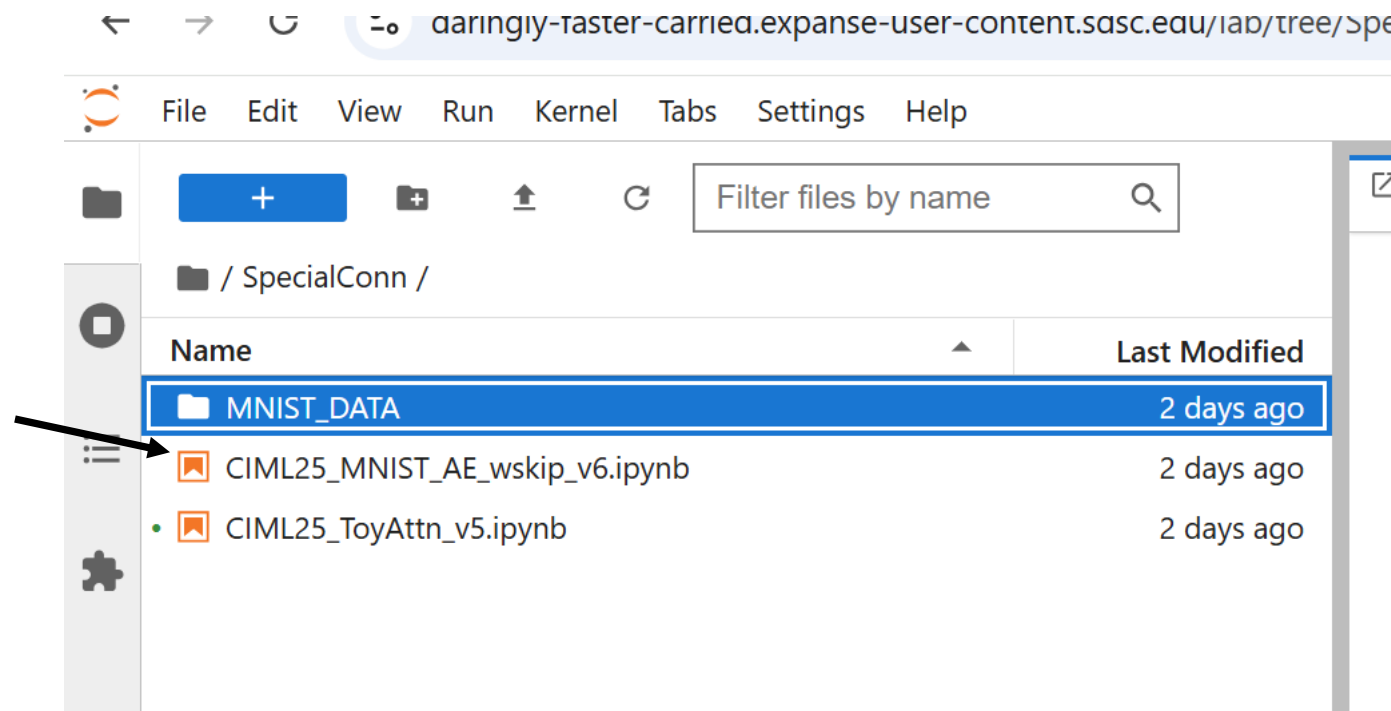
- **Review outputs to see improvements**

Login to expanse and start a notebook on
gpu-shared queue

```
$ jupyter-gpu-shared-pylight
```

In jupyter notebook session open the
(SI25) MNIST_Autoencoder_wskip
notebook

Follow instructions in the notebook



Quick overview of code

```
# -----
class MyEncoder(torch.nn.Module):
    def __init__(self):
        super(MyEncoder, self).__init__()
        #convolution layer then max pool to downsize
```

Encoder object

```
x2 = F.relu(x2)
x2 = self.max_pool_2(x2)
#print('MYINFO enc fwd, after max2, x shape:',x2.shape)
#return x2 #or x1,x2 to use skip connections
return x1,x2 #or x1,x2 to use skip connections
```

Enocde forward function returns
intermediate and last layer

```
# -----
class MyDecoder(torch.nn.Module):
    def __init__(self):
        super(MyDecoder, self).__init__()
        #convolution layer then max pool to downsize,
        self.conv1 = torch.nn.Conv2d(numfilt*2, numfilt, 3, 1, padding=1)

        #if no skip connection use in channels = numfilt for Conv2
        self.conv2 = torch.nn.Conv2d(numfilt, numfilt, 3, 1, padding=1)

        # <<<<<<<----- uncomment this, comment out the above
        #for skip connection going into conv2 use numfilt*2
        #self.conv2 = torch.nn.Conv2d(numfilt*2, numfilt, 3, 1, padding=1)
```

Decoder object

(for adding skip connections use the other `self.conv2` statement by adding/deleting comments)

Quick overview of code

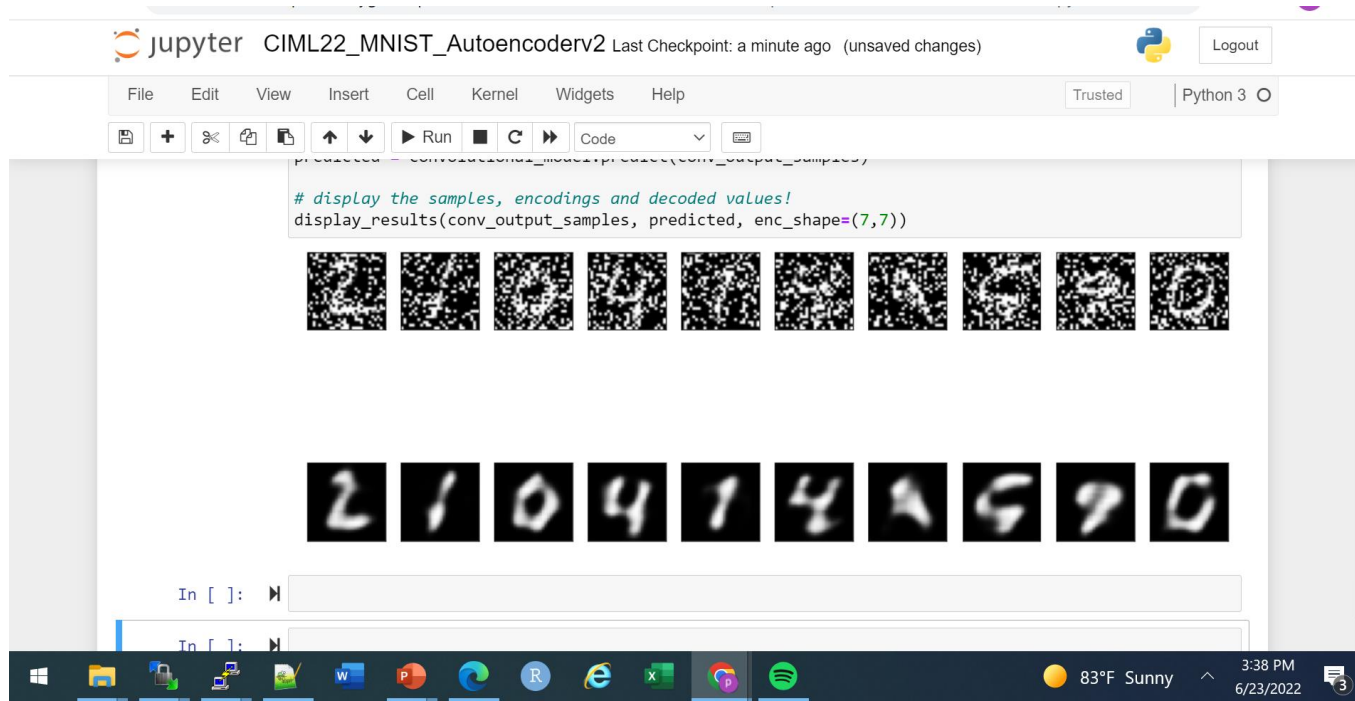
```
def forward(self, encx1,x): #or use x1,x2 inputs
    x1 = self.conv1(x)
    x1 = F.relu(x1)
    x1 = torch.nn.functional.interpolate(x1,size=(14,14),mode='nearest')
    #print('MYINFO dec fwd, after inter1, x shape:',x1.shape, 'encx1 shape:',encx1.shape)
    skip_concat_1 = torch.cat((x1,encx1), dim=1)
    #print('MYINFO, dec fwd, after concat1',skip_concat_1.shape)

    #<<<<<<----- choose if x2 should use x1 alone, or x1 concat
    x2 = self.conv2(x1)
    #x2 = self.conv2(skip_concat_1)

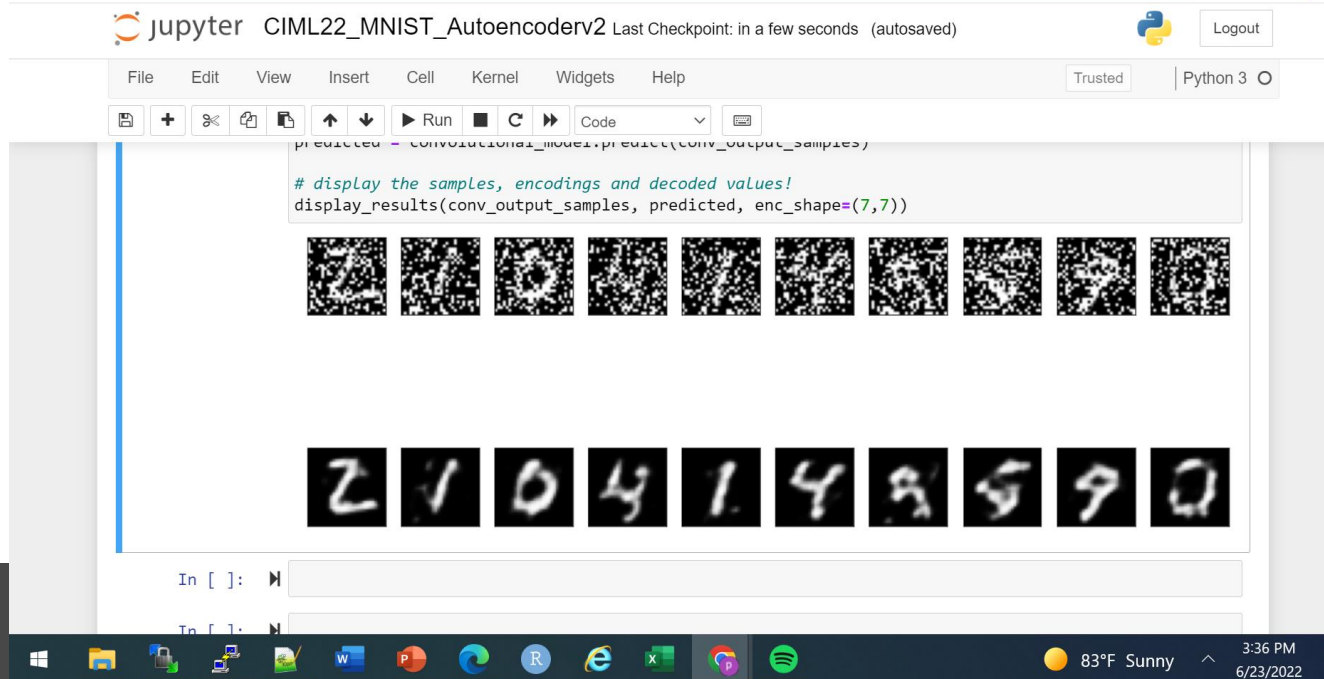
    x2 = F.relu(x2)
    x2 = torch.nn.functional.interpolate(x2,size=(28,28),mode='nearest')
    #print('MYINFO dec fwd, after inter2, x shape:',x2.shape)
    x3 = self.conv3(x2)
    x3 = F.sigmoid(x3)
    return x3
```

Decoder forward function has an argument for the first encoding layer and last encoding output

(for adding skip connections use self.conv2 with the concatenated input)



With out skip
20 epochs
Loss 0.1664



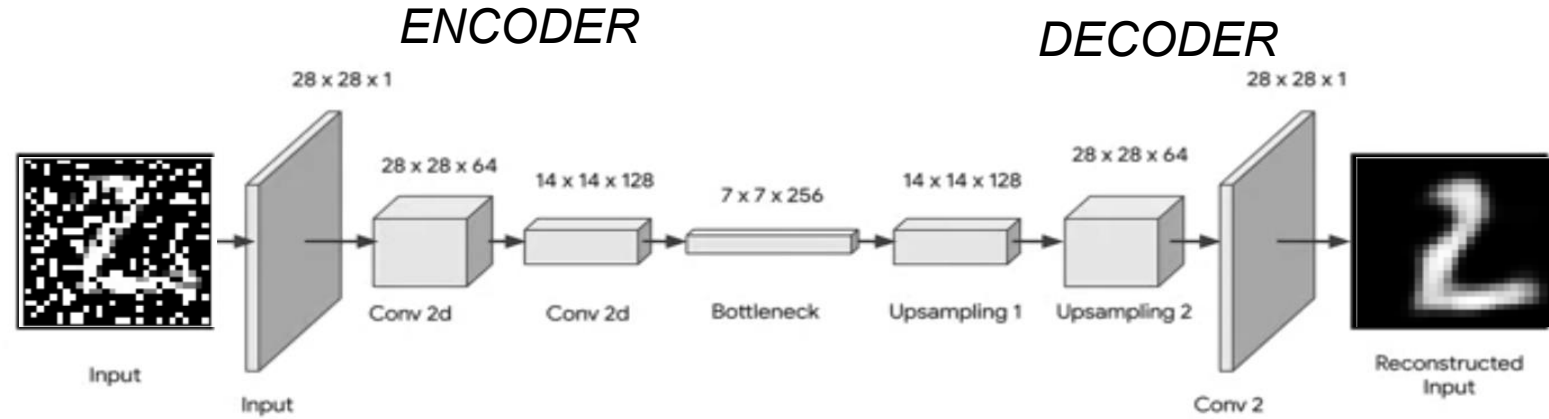
With skip,
20 epochs loss 0.14

Are the numbers a little bit
more reconstructed?


Autoencoding with Stable Diffusion

- Let's introduce the concepts and intuition behind stable diffusion


In principle, our denoising autoencoder removed noise pixels and/or filled in number pixels



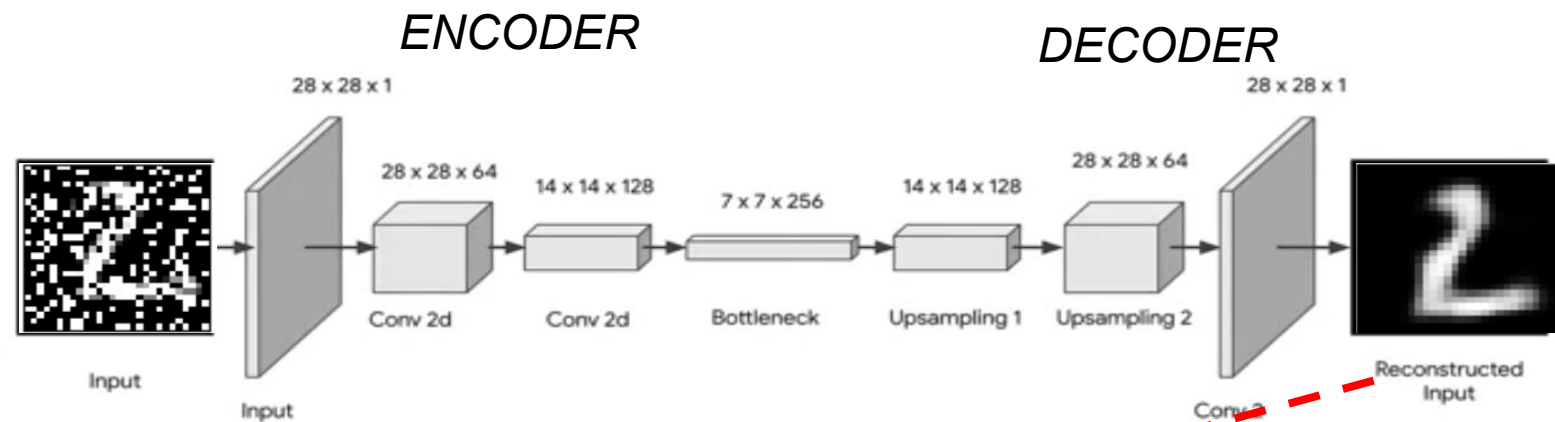
In principle, our denoising autoencoder removed noise pixels and/or filled in missing pixels



What would happen if we fed the denoised output back into the autoencoder?

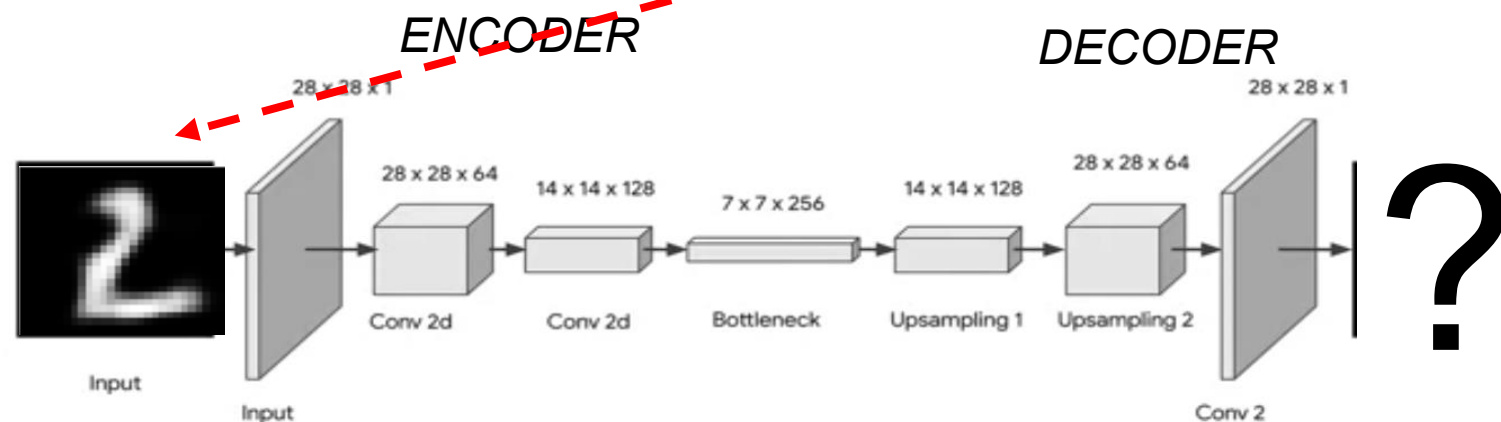


In principle, our denoising autoencoder removed noise pixels and/or filled in missing pixels




What would happen if we fed the denoised output back into the autoencoder?


- A: better reconstruction
- B: all pixels would be removed
- C: all pixels would be filled in




First step of denoising



first step of denoising




1 more step of denoising




Is it better?

First step of denoising




3 more steps of denoising



Frist step of denoising



3 more steps of denoising



Let's make this more stable, by training a network to just remove a little noise. It is like training to predict noise diffusion.

Stable Diffusion for Image Reconstruction

Jonathan Ho, Ajay Jain, and Pieter Abbeel. Denoising diffusion probabilistic models. 2020.

- **Concept:**

create a sequence of images with noise, $t=1\dots T$



Stable Diffusion for Image Reconstruction

Jonathan Ho, Ajay Jain, and Pieter Abbeel. Denoising diffusion probabilistic models. 2020.


- **Concept:**

create a sequence of images with noise, $t=1\dots T$



train the network to reconstruct image $t-1$ from image t

Note: this example is in pixel space, but it is often applied in embedding space



- From Ho et al. 2020

Early denoising steps add overall structure
Later denoising steps add more detail



- end